



University
of Windsor

Course Name

Neural Networks and Deep Learning (COMP-8610)

Document Type

Assignment 3

Professor

Dr. Alioune Ngom

Team - Members	Student ID	Contribution
Manjinder Singh	110097177	33.33%
Harbhajan Singh	110100089	33.33%
Navjot Makkar	110100116	33.33%

NN&DL_Assignment_3

June 11, 2023

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[ ]: cd /content/drive/MyDrive/Assignment_3
```

/content/drive/MyDrive/Assignment_3

```
[ ]: !ls
```

Assignment_3.ipynb	MNIST_Dataset	train-images-idx3-ubyte.gz
Assignment_3_Main.ipynb	t10k-images-idx3-ubyte.gz	train-labels-idx1-ubyte.gz
Assignment_3_old.ipynb	t10k-labels-idx1-ubyte.gz	

1 Que 1:

Download the benchmark dataset, MNIST, from <http://yann.lecun.com/exdb/mnist/>. Implement multi-class classification for recognizing handwritten digits (also known as multiclass logistic regression —this is simply a feedforward neural network with k output neurons, with one output neuron for each class, and each output neuron o_i returns the probability that the input datapoint x_j is in class i) and try it on MNIST.

Comments: No need to implement almost anything in DL by your own (this is true in general); the software framework (ie, the DL platform) typically provides implementations for all the things discussed in class, such as the learning algorithms, the regularizations methods, the cross-validation methods, etc.

Use your favorite deep learning platform. A few candidates: 1. Marvin from <http://marvin.is/> 2. Caffe from <http://caffe.berkeleyvision.org> 3. TensorFlow from <https://www.tensorflow.org> 4. Pylearn2 from <http://deeplearning.net/software/pylearn2/> 5. Theano, Torch, Lasagne, etc. See more platforms at http://deeplearning.net/software_links/.

Read the tutorial about your selected platform (eg, for TensorFlow: <https://www.tensorflow.org/tutorials>), try it on MNIST; note that the first few examples in the tutorials are typically on MNIST or other simple image datasets, so you can follow the steps.

Comments: MNIST is a standard dataset for machine learning and also deep learning. It's good to try it on one shallow neural network (with one output neuron; eg, for recognizing a character A from a not-A character) before trying it on a deep neural network with multiple outputs. Downloading

the dataset from other places in preprocessed format is allowed, but practicing how to read the dataset prepares you for other new datasets you may be interested in (thus, please, read the MNIST website carefully).

1.1 Que 1: Part 1

1. Try the basic minibatch SGD as your learning algorithm. It is recommended to try different initializations, different batch sizes, and different learning rates, in order to get a sense about how to tune the hyperparameters (batch size, and, learning rate). Remember to create and use validation dataset!. it will be very useful for you to read Chapter-11 of the textbook.

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
import urllib.request
import gzip
from google.colab import drive
import os

# Step 1: Download and read the MNIST dataset
def download_mnist():
    base_url = 'http://yann.lecun.com/exdb/mnist/'
    file_names = ['train-images-idx3-ubyte.gz', 'train-labels-idx1-ubyte.gz',
                  't10k-images-idx3-ubyte.gz', 't10k-labels-idx1-ubyte.gz']

    for file_name in file_names:
        urllib.request.urlretrieve(base_url + file_name, file_name)
```

```
[ ]: def read_mnist_images(filename):
    with gzip.open(filename, 'rb') as f:
        data = np.frombuffer(f.read(), np.uint8, offset=16)
    return data.reshape(-1, 784)

def read_mnist_labels(filename):
    with gzip.open(filename, 'rb') as f:
        data = np.frombuffer(f.read(), np.uint8, offset=8)
    return data
```

```
[ ]: # Create MNIST folder in Google Drive
!mkdir -p '/content/drive/MyDrive/Assignment_3/MNIST_Dataset'
```

```
[ ]: # Set the Google Drive folder as the working directory
os.chdir('/content/drive/MyDrive/Assignment_3/MNIST_Dataset')
```

```
[ ]: # Download MNIST dataset
download_mnist()
```

```
# Read training and testing images and labels
train_images = read_mnist_images('train-images-idx3-ubyte.gz')
train_labels = read_mnist_labels('train-labels-idx1-ubyte.gz')
test_images = read_mnist_images('t10k-images-idx3-ubyte.gz')
test_labels = read_mnist_labels('t10k-labels-idx1-ubyte.gz')
```

```
[ ]: print("Size of Training Images Dataset : ", len(train_images))
print("Size of Training Images Dataset : ", len(test_labels))
```

```
Size of Training Images Dataset : 60000
Size of Training Images Dataset : 10000
```

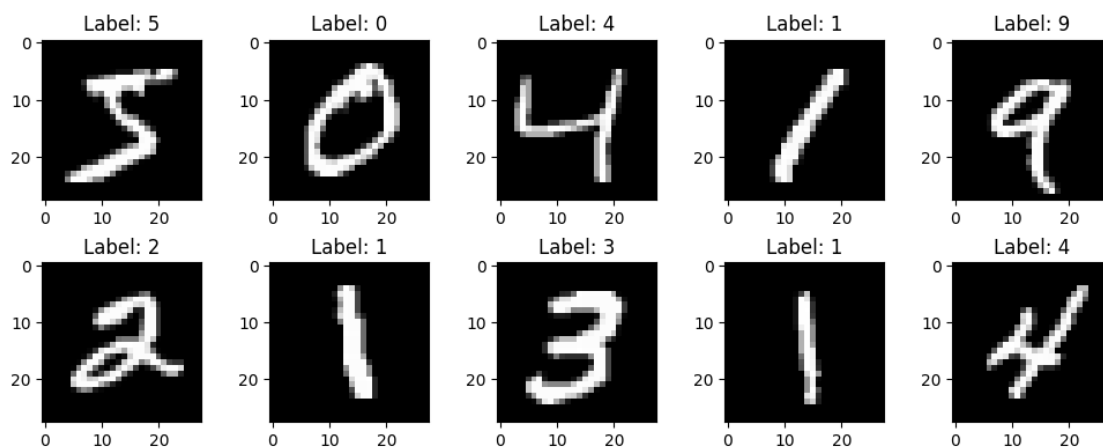
```
[ ]: train_images[1].shape
```

```
[ ]: (784,)
```

```
[ ]: # Display some training images
fig, axs = plt.subplots(2, 5, figsize=(10, 4))
axs = axs.flatten()

for i in range(10):
    img = train_images[i].reshape(28, 28)
    axs[i].imshow(img, cmap='gray')
    axs[i].set_title(f"Label: {train_labels[i]}")

plt.tight_layout()
plt.show()
```




1.1.1 Data Preprocessing

```
[ ]: # Step 2: Data preprocessing
train_images = train_images / 255.0
test_images = test_images / 255.0

# Showing images after preprocessing
fig, axs = plt.subplots(1, 5, figsize=(10, 4))
axs = axs.flatten()

for i in range(5):
    img = train_images[i].reshape(28, 28)
    axs[i].imshow(img, cmap='gray')
    axs[i].set_title(f"Label: {train_labels[i]}")

plt.tight_layout()
plt.show()
```



```
[ ]: # Step 2: Prepare the data
num_classes = 10
num_features = train_images.shape[1]
num_samples = train_images.shape[0]
```

```
[ ]: import tensorflow as tf
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.wrappers.scikit_learn import KerasClassifier
```

```
[23]: # Preprocess the MNIST dataset
input_dim = train_images.shape[1]
x_train = train_images.reshape(-1, 784)
x_test = test_images.reshape(-1, 784)
y_train = tf.keras.utils.to_categorical(train_labels, num_classes=10)
y_test = tf.keras.utils.to_categorical(test_labels, num_classes=10)
```

```
# Create a validation dataset
x_train, x_val, y_train, y_val = train_test_split(x_train, y_train, test_size=0.
↳1, random_state=42)
```

```
[ ]: # Define a function to create your model with the desired architecture
def create_model(learning_rate, hidden_units):
    model = Sequential()
    model.add(Dense(hidden_units, activation='relu', input_shape=(input_dim,)))
    model.add(Dense(num_classes, activation='sigmoid'))
    optimizer = SGD(learning_rate=learning_rate)
    model.compile(loss='categorical_crossentropy', optimizer=optimizer,
↳metrics=['accuracy'])
    return model

# Create the KerasClassifier wrapper
model = KerasClassifier(build_fn=create_model, epochs=11, batch_size=64,
↳verbose=1)

# Define the hyperparameters grid
param_grid = {
    'learning_rate': [0.001, 0.01, 0.1],
    'hidden_units': [32, 64, 128]
}

# Perform grid search cross-validation
grid = GridSearchCV(estimator=model, param_grid=param_grid, cv=5)
grid_result = grid.fit(x_train, y_train)

print("\n\n")
# Print the best hyperparameters and accuracy
print("Best Hyperparameters: ", grid_result.best_params_)
print("Best Accuracy: ", grid_result.best_score_)
```

```
<ipython-input-16-bb1ba6c3502e>:11: DeprecationWarning: KerasClassifier is
deprecated, use Sci-Keras (https://github.com/adriangb/scikeras) instead. See
https://www.adriangb.com/scikeras/stable/migration.html for help migrating.
    model = KerasClassifier(build_fn=create_model, epochs=11, batch_size=64,
verbose=1)
```

```
Epoch 1/11
```

```
675/675 [=====] - 2s 2ms/step - loss: 2.1072 -
accuracy: 0.2839
```

```
Epoch 2/11
```

```
675/675 [=====] - 1s 2ms/step - loss: 1.7024 -
accuracy: 0.6174
```

```
Epoch 3/11
```

```
675/675 [=====] - 1s 2ms/step - loss: 1.3518 -
```

```

accuracy: 0.7499
Epoch 4/11
675/675 [=====] - 1s 2ms/step - loss: 1.0960 -
accuracy: 0.7933
Epoch 5/11
675/675 [=====] - 1s 2ms/step - loss: 0.9209 -
accuracy: 0.8138
Epoch 6/11
675/675 [=====] - 1s 2ms/step - loss: 0.8023 -
accuracy: 0.8291
Epoch 7/11
675/675 [=====] - 1s 2ms/step - loss: 0.7202 -
accuracy: 0.8399
Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.6610 -
accuracy: 0.8481
Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.6162 -
accuracy: 0.8534
Epoch 10/11
675/675 [=====] - 2s 3ms/step - loss: 0.5813 -
accuracy: 0.8585
Epoch 11/11
675/675 [=====] - 1s 2ms/step - loss: 0.5533 -
accuracy: 0.8629
169/169 [=====] - 0s 2ms/step - loss: 0.5447 -
accuracy: 0.8634
Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 2.0535 -
accuracy: 0.3253
Epoch 2/11
675/675 [=====] - 1s 2ms/step - loss: 1.5785 -
accuracy: 0.5901
Epoch 3/11
675/675 [=====] - 2s 2ms/step - loss: 1.2628 -
accuracy: 0.6912
Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 1.0538 -
accuracy: 0.7501
Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.9116 -
accuracy: 0.7856
Epoch 6/11
675/675 [=====] - 1s 2ms/step - loss: 0.8113 -
accuracy: 0.8069
Epoch 7/11
675/675 [=====] - 1s 2ms/step - loss: 0.7379 -
accuracy: 0.8214

```

Epoch 8/11
675/675 [=====] - 1s 2ms/step - loss: 0.6823 -
accuracy: 0.8322

Epoch 9/11
675/675 [=====] - 1s 2ms/step - loss: 0.6387 -
accuracy: 0.8400

Epoch 10/11
675/675 [=====] - 1s 2ms/step - loss: 0.6038 -
accuracy: 0.8476

Epoch 11/11
675/675 [=====] - 1s 2ms/step - loss: 0.5749 -
accuracy: 0.8536

169/169 [=====] - 1s 2ms/step - loss: 0.5706 -
accuracy: 0.8571

Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 2.1303 -
accuracy: 0.2753

Epoch 2/11
675/675 [=====] - 1s 2ms/step - loss: 1.7394 -
accuracy: 0.5300

Epoch 3/11
675/675 [=====] - 1s 2ms/step - loss: 1.4137 -
accuracy: 0.6708

Epoch 4/11
675/675 [=====] - 1s 2ms/step - loss: 1.1685 -
accuracy: 0.7523

Epoch 5/11
675/675 [=====] - 1s 2ms/step - loss: 0.9916 -
accuracy: 0.7916

Epoch 6/11
675/675 [=====] - 2s 3ms/step - loss: 0.8661 -
accuracy: 0.8117

Epoch 7/11
675/675 [=====] - 3s 4ms/step - loss: 0.7760 -
accuracy: 0.8258

Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.7094 -
accuracy: 0.8347

Epoch 9/11
675/675 [=====] - 2s 2ms/step - loss: 0.6586 -
accuracy: 0.8411

Epoch 10/11
675/675 [=====] - 1s 2ms/step - loss: 0.6189 -
accuracy: 0.8471

Epoch 11/11
675/675 [=====] - 1s 2ms/step - loss: 0.5869 -
accuracy: 0.8518

169/169 [=====] - 0s 2ms/step - loss: 0.5693 -


```

accuracy: 0.8555
Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 2.1369 -
accuracy: 0.2748
Epoch 2/11
675/675 [=====] - 2s 2ms/step - loss: 1.7316 -
accuracy: 0.5704
Epoch 3/11
675/675 [=====] - 2s 3ms/step - loss: 1.4097 -
accuracy: 0.6673
Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 1.1790 -
accuracy: 0.7269
Epoch 5/11
675/675 [=====] - 1s 2ms/step - loss: 1.0154 -
accuracy: 0.7686
Epoch 6/11
675/675 [=====] - 1s 2ms/step - loss: 0.8970 -
accuracy: 0.7944
Epoch 7/11
675/675 [=====] - 1s 2ms/step - loss: 0.8089 -
accuracy: 0.8119
Epoch 8/11
675/675 [=====] - 1s 2ms/step - loss: 0.7416 -
accuracy: 0.8248
Epoch 9/11
675/675 [=====] - 1s 2ms/step - loss: 0.6889 -
accuracy: 0.8348
Epoch 10/11
675/675 [=====] - 1s 2ms/step - loss: 0.6465 -
accuracy: 0.8427
Epoch 11/11
675/675 [=====] - 1s 2ms/step - loss: 0.6118 -
accuracy: 0.8490
169/169 [=====] - 0s 2ms/step - loss: 0.6071 -
accuracy: 0.8476
Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 2.1650 -
accuracy: 0.2593
Epoch 2/11
675/675 [=====] - 1s 2ms/step - loss: 1.8125 -
accuracy: 0.5514
Epoch 3/11
675/675 [=====] - 1s 2ms/step - loss: 1.4598 -
accuracy: 0.7019
Epoch 4/11
675/675 [=====] - 1s 2ms/step - loss: 1.1853 -
accuracy: 0.7581

```

Epoch 5/11
675/675 [=====] - 1s 2ms/step - loss: 0.9956 -
accuracy: 0.7913

Epoch 6/11
675/675 [=====] - 1s 2ms/step - loss: 0.8648 -
accuracy: 0.8148

Epoch 7/11
675/675 [=====] - 2s 3ms/step - loss: 0.7720 -
accuracy: 0.8294

Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.7037 -
accuracy: 0.8404

Epoch 9/11
675/675 [=====] - 1s 2ms/step - loss: 0.6516 -
accuracy: 0.8483

Epoch 10/11
675/675 [=====] - 1s 2ms/step - loss: 0.6106 -
accuracy: 0.8549

Epoch 11/11
675/675 [=====] - 1s 2ms/step - loss: 0.5777 -
accuracy: 0.8607

169/169 [=====] - 0s 2ms/step - loss: 0.5555 -
accuracy: 0.8668

Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 1.1842 -
accuracy: 0.6800

Epoch 2/11
675/675 [=====] - 1s 2ms/step - loss: 0.5282 -
accuracy: 0.8640

Epoch 3/11
675/675 [=====] - 1s 2ms/step - loss: 0.4213 -
accuracy: 0.8867

Epoch 4/11
675/675 [=====] - 2s 2ms/step - loss: 0.3756 -
accuracy: 0.8958

Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.3481 -
accuracy: 0.9025

Epoch 6/11
675/675 [=====] - 2s 3ms/step - loss: 0.3293 -
accuracy: 0.9067

Epoch 7/11
675/675 [=====] - 1s 2ms/step - loss: 0.3151 -
accuracy: 0.9104

Epoch 8/11
675/675 [=====] - 1s 2ms/step - loss: 0.3030 -
accuracy: 0.9139

Epoch 9/11

```

675/675 [=====] - 1s 2ms/step - loss: 0.2928 -
accuracy: 0.9167
Epoch 10/11
675/675 [=====] - 1s 2ms/step - loss: 0.2840 -
accuracy: 0.9202
Epoch 11/11
675/675 [=====] - 1s 2ms/step - loss: 0.2757 -
accuracy: 0.9223
169/169 [=====] - 0s 2ms/step - loss: 0.2811 -
accuracy: 0.9205
Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 1.2126 -
accuracy: 0.6791
Epoch 2/11
675/675 [=====] - 2s 3ms/step - loss: 0.5228 -
accuracy: 0.8691
Epoch 3/11
675/675 [=====] - 2s 3ms/step - loss: 0.4155 -
accuracy: 0.8864
Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 0.3712 -
accuracy: 0.8968
Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.3447 -
accuracy: 0.9032
Epoch 6/11
675/675 [=====] - 2s 2ms/step - loss: 0.3260 -
accuracy: 0.9084
Epoch 7/11
675/675 [=====] - 1s 2ms/step - loss: 0.3115 -
accuracy: 0.9118
Epoch 8/11
675/675 [=====] - 1s 2ms/step - loss: 0.2996 -
accuracy: 0.9152
Epoch 9/11
675/675 [=====] - 1s 2ms/step - loss: 0.2892 -
accuracy: 0.9183
Epoch 10/11
675/675 [=====] - 1s 2ms/step - loss: 0.2805 -
accuracy: 0.9206
Epoch 11/11
675/675 [=====] - 1s 2ms/step - loss: 0.2726 -
accuracy: 0.9225
169/169 [=====] - 0s 2ms/step - loss: 0.2814 -
accuracy: 0.9215
Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 1.0690 -
accuracy: 0.7380

```

Epoch 2/11
675/675 [=====] - 2s 3ms/step - loss: 0.4929 -
accuracy: 0.8737

Epoch 3/11
675/675 [=====] - 2s 3ms/step - loss: 0.4031 -
accuracy: 0.8908

Epoch 4/11
675/675 [=====] - 1s 2ms/step - loss: 0.3630 -
accuracy: 0.9010

Epoch 5/11
675/675 [=====] - 1s 2ms/step - loss: 0.3381 -
accuracy: 0.9066

Epoch 6/11
675/675 [=====] - 1s 2ms/step - loss: 0.3200 -
accuracy: 0.9114

Epoch 7/11
675/675 [=====] - 1s 2ms/step - loss: 0.3061 -
accuracy: 0.9147

Epoch 8/11
675/675 [=====] - 1s 2ms/step - loss: 0.2944 -
accuracy: 0.9179

Epoch 9/11
675/675 [=====] - 1s 2ms/step - loss: 0.2844 -
accuracy: 0.9207

Epoch 10/11
675/675 [=====] - 1s 2ms/step - loss: 0.2758 -
accuracy: 0.9230

Epoch 11/11
675/675 [=====] - 2s 2ms/step - loss: 0.2682 -
accuracy: 0.9248

169/169 [=====] - 0s 2ms/step - loss: 0.2661 -
accuracy: 0.9254

Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 1.2580 -
accuracy: 0.6683

Epoch 2/11
675/675 [=====] - 1s 2ms/step - loss: 0.5407 -
accuracy: 0.8632

Epoch 3/11
675/675 [=====] - 1s 2ms/step - loss: 0.4251 -
accuracy: 0.8852

Epoch 4/11
675/675 [=====] - 1s 2ms/step - loss: 0.3778 -
accuracy: 0.8945

Epoch 5/11
675/675 [=====] - 1s 2ms/step - loss: 0.3498 -
accuracy: 0.9007

Epoch 6/11

675/675 [=====] - 2s 3ms/step - loss: 0.3298 -
accuracy: 0.9056
Epoch 7/11
675/675 [=====] - 2s 3ms/step - loss: 0.3141 -
accuracy: 0.9098
Epoch 8/11
675/675 [=====] - 2s 2ms/step - loss: 0.3012 -
accuracy: 0.9136
Epoch 9/11
675/675 [=====] - 1s 2ms/step - loss: 0.2901 -
accuracy: 0.9163
Epoch 10/11
675/675 [=====] - 1s 2ms/step - loss: 0.2807 -
accuracy: 0.9195
Epoch 11/11
675/675 [=====] - 1s 2ms/step - loss: 0.2723 -
accuracy: 0.9220
169/169 [=====] - 0s 2ms/step - loss: 0.2904 -
accuracy: 0.9179
Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 1.1019 -
accuracy: 0.7120
Epoch 2/11
675/675 [=====] - 1s 2ms/step - loss: 0.5257 -
accuracy: 0.8631
Epoch 3/11
675/675 [=====] - 1s 2ms/step - loss: 0.4178 -
accuracy: 0.8878
Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 0.3699 -
accuracy: 0.8985
Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.3416 -
accuracy: 0.9053
Epoch 6/11
675/675 [=====] - 2s 3ms/step - loss: 0.3220 -
accuracy: 0.9102
Epoch 7/11
675/675 [=====] - 1s 2ms/step - loss: 0.3067 -
accuracy: 0.9145
Epoch 8/11
675/675 [=====] - 1s 2ms/step - loss: 0.2945 -
accuracy: 0.9178
Epoch 9/11
675/675 [=====] - 1s 2ms/step - loss: 0.2842 -
accuracy: 0.9197
Epoch 10/11
675/675 [=====] - 1s 2ms/step - loss: 0.2750 -

```

accuracy: 0.9220
Epoch 11/11
675/675 [=====] - 1s 2ms/step - loss: 0.2668 -
accuracy: 0.9251
169/169 [=====] - 1s 2ms/step - loss: 0.2668 -
accuracy: 0.9255
Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 0.4600 -
accuracy: 0.8703
Epoch 2/11
675/675 [=====] - 1s 2ms/step - loss: 0.2712 -
accuracy: 0.9222
Epoch 3/11
675/675 [=====] - 1s 2ms/step - loss: 0.2270 -
accuracy: 0.9353
Epoch 4/11
675/675 [=====] - 1s 2ms/step - loss: 0.1969 -
accuracy: 0.9434
Epoch 5/11
675/675 [=====] - 1s 2ms/step - loss: 0.1744 -
accuracy: 0.9505
Epoch 6/11
675/675 [=====] - 1s 2ms/step - loss: 0.1576 -
accuracy: 0.9563
Epoch 7/11
675/675 [=====] - 1s 2ms/step - loss: 0.1444 -
accuracy: 0.9588
Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.1328 -
accuracy: 0.9626
Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.1233 -
accuracy: 0.9647
Epoch 10/11
675/675 [=====] - 2s 3ms/step - loss: 0.1153 -
accuracy: 0.9672
Epoch 11/11
675/675 [=====] - 1s 2ms/step - loss: 0.1085 -
accuracy: 0.9691
169/169 [=====] - 0s 2ms/step - loss: 0.1434 -
accuracy: 0.9592
Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 0.4496 -
accuracy: 0.8743
Epoch 2/11
675/675 [=====] - 1s 2ms/step - loss: 0.2631 -
accuracy: 0.9235
Epoch 3/11

```

675/675 [=====] - 2s 3ms/step - loss: 0.2161 -
accuracy: 0.9389
Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 0.1876 -
accuracy: 0.9467
Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.1665 -
accuracy: 0.9520
Epoch 6/11
675/675 [=====] - 1s 2ms/step - loss: 0.1518 -
accuracy: 0.9550
Epoch 7/11
675/675 [=====] - 1s 2ms/step - loss: 0.1400 -
accuracy: 0.9587
Epoch 8/11
675/675 [=====] - 1s 2ms/step - loss: 0.1299 -
accuracy: 0.9618
Epoch 9/11
675/675 [=====] - 1s 2ms/step - loss: 0.1209 -
accuracy: 0.9638
Epoch 10/11
675/675 [=====] - 1s 2ms/step - loss: 0.1143 -
accuracy: 0.9665
Epoch 11/11
675/675 [=====] - 1s 2ms/step - loss: 0.1075 -
accuracy: 0.9683
169/169 [=====] - 0s 2ms/step - loss: 0.1442 -
accuracy: 0.9581
Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 0.4427 -
accuracy: 0.8773
Epoch 2/11
675/675 [=====] - 1s 2ms/step - loss: 0.2520 -
accuracy: 0.9278
Epoch 3/11
675/675 [=====] - 1s 2ms/step - loss: 0.2049 -
accuracy: 0.9406
Epoch 4/11
675/675 [=====] - 1s 2ms/step - loss: 0.1752 -
accuracy: 0.9499
Epoch 5/11
675/675 [=====] - 1s 2ms/step - loss: 0.1563 -
accuracy: 0.9547
Epoch 6/11
675/675 [=====] - 1s 2ms/step - loss: 0.1415 -
accuracy: 0.9587
Epoch 7/11
675/675 [=====] - 2s 3ms/step - loss: 0.1298 -

```

accuracy: 0.9619
Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.1201 -
accuracy: 0.9652
Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.1112 -
accuracy: 0.9681
Epoch 10/11
675/675 [=====] - 1s 2ms/step - loss: 0.1038 -
accuracy: 0.9697
Epoch 11/11
675/675 [=====] - 1s 2ms/step - loss: 0.0983 -
accuracy: 0.9713
169/169 [=====] - 0s 2ms/step - loss: 0.1370 -
accuracy: 0.9611
Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 0.4515 -
accuracy: 0.8738
Epoch 2/11
675/675 [=====] - 2s 3ms/step - loss: 0.2648 -
accuracy: 0.9231
Epoch 3/11
675/675 [=====] - 2s 3ms/step - loss: 0.2195 -
accuracy: 0.9366
Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 0.1905 -
accuracy: 0.9454
Epoch 5/11
675/675 [=====] - 1s 2ms/step - loss: 0.1689 -
accuracy: 0.9515
Epoch 6/11
675/675 [=====] - 1s 2ms/step - loss: 0.1538 -
accuracy: 0.9557
Epoch 7/11
675/675 [=====] - 1s 2ms/step - loss: 0.1405 -
accuracy: 0.9584
Epoch 8/11
675/675 [=====] - 1s 2ms/step - loss: 0.1292 -
accuracy: 0.9628
Epoch 9/11
675/675 [=====] - 1s 2ms/step - loss: 0.1206 -
accuracy: 0.9648
Epoch 10/11
675/675 [=====] - 1s 2ms/step - loss: 0.1134 -
accuracy: 0.9669
Epoch 11/11
675/675 [=====] - 1s 2ms/step - loss: 0.1072 -
accuracy: 0.9681

```



```

169/169 [=====] - 0s 2ms/step - loss: 0.1490 -
accuracy: 0.9568
Epoch 1/11
675/675 [=====] - 2s 3ms/step - loss: 0.4585 -
accuracy: 0.8710
Epoch 2/11
675/675 [=====] - 2s 3ms/step - loss: 0.2614 -
accuracy: 0.9259
Epoch 3/11
675/675 [=====] - 1s 2ms/step - loss: 0.2129 -
accuracy: 0.9393
Epoch 4/11
675/675 [=====] - 1s 2ms/step - loss: 0.1827 -
accuracy: 0.9475
Epoch 5/11
675/675 [=====] - 2s 2ms/step - loss: 0.1609 -
accuracy: 0.9539
Epoch 6/11
675/675 [=====] - 2s 3ms/step - loss: 0.1441 -
accuracy: 0.9591
Epoch 7/11
675/675 [=====] - 2s 3ms/step - loss: 0.1319 -
accuracy: 0.9623
Epoch 8/11
675/675 [=====] - 1s 2ms/step - loss: 0.1218 -
accuracy: 0.9656
Epoch 9/11
675/675 [=====] - 1s 2ms/step - loss: 0.1135 -
accuracy: 0.9671
Epoch 10/11
675/675 [=====] - 1s 2ms/step - loss: 0.1071 -
accuracy: 0.9688
Epoch 11/11
675/675 [=====] - 1s 2ms/step - loss: 0.1010 -
accuracy: 0.9707
169/169 [=====] - 0s 2ms/step - loss: 0.1297 -
accuracy: 0.9608
Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 2.0635 -
accuracy: 0.3171
Epoch 2/11
675/675 [=====] - 2s 2ms/step - loss: 1.6371 -
accuracy: 0.5908
Epoch 3/11
675/675 [=====] - 2s 3ms/step - loss: 1.3018 -
accuracy: 0.7126
Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 1.0656 -

```

```

accuracy: 0.7716
Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.9068 -
accuracy: 0.8036
Epoch 6/11
675/675 [=====] - 2s 2ms/step - loss: 0.7983 -
accuracy: 0.8234
Epoch 7/11
675/675 [=====] - 2s 2ms/step - loss: 0.7211 -
accuracy: 0.8367
Epoch 8/11
675/675 [=====] - 2s 2ms/step - loss: 0.6639 -
accuracy: 0.8464
Epoch 9/11
675/675 [=====] - 2s 2ms/step - loss: 0.6199 -
accuracy: 0.8531
Epoch 10/11
675/675 [=====] - 2s 2ms/step - loss: 0.5851 -
accuracy: 0.8581
Epoch 11/11
675/675 [=====] - 2s 2ms/step - loss: 0.5569 -
accuracy: 0.8628
169/169 [=====] - 1s 2ms/step - loss: 0.5456 -
accuracy: 0.8640
Epoch 1/11
675/675 [=====] - 3s 3ms/step - loss: 2.0986 -
accuracy: 0.3148
Epoch 2/11
675/675 [=====] - 2s 2ms/step - loss: 1.6424 -
accuracy: 0.6170
Epoch 3/11
675/675 [=====] - 1s 2ms/step - loss: 1.3034 -
accuracy: 0.7211
Epoch 4/11
675/675 [=====] - 2s 2ms/step - loss: 1.0674 -
accuracy: 0.7730
Epoch 5/11
675/675 [=====] - 1s 2ms/step - loss: 0.9092 -
accuracy: 0.8003
Epoch 6/11
675/675 [=====] - 2s 2ms/step - loss: 0.8007 -
accuracy: 0.8178
Epoch 7/11
675/675 [=====] - 2s 2ms/step - loss: 0.7230 -
accuracy: 0.8305
Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.6652 -
accuracy: 0.8403

```

Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.6206 -
accuracy: 0.8480

Epoch 10/11
675/675 [=====] - 2s 3ms/step - loss: 0.5852 -
accuracy: 0.8538

Epoch 11/11
675/675 [=====] - 1s 2ms/step - loss: 0.5566 -
accuracy: 0.8596
169/169 [=====] - 0s 2ms/step - loss: 0.5543 -
accuracy: 0.8631

Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 2.0718 -
accuracy: 0.3389

Epoch 2/11
675/675 [=====] - 2s 2ms/step - loss: 1.6336 -
accuracy: 0.6017

Epoch 3/11
675/675 [=====] - 1s 2ms/step - loss: 1.3072 -
accuracy: 0.7150

Epoch 4/11
675/675 [=====] - 2s 2ms/step - loss: 1.0760 -
accuracy: 0.7719

Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.9190 -
accuracy: 0.8033

Epoch 6/11
675/675 [=====] - 2s 3ms/step - loss: 0.8112 -
accuracy: 0.8202

Epoch 7/11
675/675 [=====] - 2s 2ms/step - loss: 0.7338 -
accuracy: 0.8327

Epoch 8/11
675/675 [=====] - 2s 2ms/step - loss: 0.6760 -
accuracy: 0.8417

Epoch 9/11
675/675 [=====] - 1s 2ms/step - loss: 0.6312 -
accuracy: 0.8487

Epoch 10/11
675/675 [=====] - 1s 2ms/step - loss: 0.5955 -
accuracy: 0.8543

Epoch 11/11
675/675 [=====] - 2s 2ms/step - loss: 0.5662 -
accuracy: 0.8593
169/169 [=====] - 1s 3ms/step - loss: 0.5501 -
accuracy: 0.8630

Epoch 1/11
675/675 [=====] - 3s 3ms/step - loss: 2.1824 -

```

accuracy: 0.2367
Epoch 2/11
675/675 [=====] - 2s 2ms/step - loss: 1.7334 -
accuracy: 0.6232
Epoch 3/11
675/675 [=====] - 2s 2ms/step - loss: 1.3807 -
accuracy: 0.7326
Epoch 4/11
675/675 [=====] - 2s 2ms/step - loss: 1.1291 -
accuracy: 0.7773
Epoch 5/11
675/675 [=====] - 2s 2ms/step - loss: 0.9579 -
accuracy: 0.8011
Epoch 6/11
675/675 [=====] - 2s 3ms/step - loss: 0.8409 -
accuracy: 0.8172
Epoch 7/11
675/675 [=====] - 2s 3ms/step - loss: 0.7582 -
accuracy: 0.8276
Epoch 8/11
675/675 [=====] - 3s 4ms/step - loss: 0.6970 -
accuracy: 0.8360
Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.6499 -
accuracy: 0.8435
Epoch 10/11
675/675 [=====] - 2s 3ms/step - loss: 0.6125 -
accuracy: 0.8496
Epoch 11/11
675/675 [=====] - 2s 2ms/step - loss: 0.5818 -
accuracy: 0.8549
169/169 [=====] - 0s 2ms/step - loss: 0.5808 -
accuracy: 0.8555
Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 2.0995 -
accuracy: 0.2950
Epoch 2/11
675/675 [=====] - 2s 2ms/step - loss: 1.6475 -
accuracy: 0.6154
Epoch 3/11
675/675 [=====] - 2s 2ms/step - loss: 1.3001 -
accuracy: 0.7327
Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 1.0629 -
accuracy: 0.7806
Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.9063 -
accuracy: 0.8047

```

Epoch 6/11
675/675 [=====] - 2s 3ms/step - loss: 0.7995 -
accuracy: 0.8219

Epoch 7/11
675/675 [=====] - 2s 2ms/step - loss: 0.7232 -
accuracy: 0.8340

Epoch 8/11
675/675 [=====] - 2s 2ms/step - loss: 0.6663 -
accuracy: 0.8420

Epoch 9/11
675/675 [=====] - 2s 2ms/step - loss: 0.6222 -
accuracy: 0.8500

Epoch 10/11
675/675 [=====] - 2s 2ms/step - loss: 0.5870 -
accuracy: 0.8562

Epoch 11/11
675/675 [=====] - 2s 2ms/step - loss: 0.5583 -
accuracy: 0.8609

169/169 [=====] - 1s 2ms/step - loss: 0.5406 -
accuracy: 0.8646

Epoch 1/11
675/675 [=====] - 3s 3ms/step - loss: 1.0401 -
accuracy: 0.7384

Epoch 2/11
675/675 [=====] - 2s 2ms/step - loss: 0.4832 -
accuracy: 0.8757

Epoch 3/11
675/675 [=====] - 2s 2ms/step - loss: 0.3981 -
accuracy: 0.8924

Epoch 4/11
675/675 [=====] - 2s 2ms/step - loss: 0.3595 -
accuracy: 0.9013

Epoch 5/11
675/675 [=====] - 2s 2ms/step - loss: 0.3352 -
accuracy: 0.9070

Epoch 6/11
675/675 [=====] - 2s 2ms/step - loss: 0.3169 -
accuracy: 0.9115

Epoch 7/11
675/675 [=====] - 2s 2ms/step - loss: 0.3019 -
accuracy: 0.9157

Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.2894 -
accuracy: 0.9193

Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.2782 -
accuracy: 0.9228

Epoch 10/11

```

675/675 [=====] - 2s 3ms/step - loss: 0.2686 -
accuracy: 0.9257
Epoch 11/11
675/675 [=====] - 2s 2ms/step - loss: 0.2597 -
accuracy: 0.9283
169/169 [=====] - 0s 2ms/step - loss: 0.2633 -
accuracy: 0.9277
Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 1.0434 -
accuracy: 0.7363
Epoch 2/11
675/675 [=====] - 2s 3ms/step - loss: 0.4906 -
accuracy: 0.8730
Epoch 3/11
675/675 [=====] - 2s 2ms/step - loss: 0.3979 -
accuracy: 0.8920
Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 0.3560 -
accuracy: 0.9017
Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.3302 -
accuracy: 0.9077
Epoch 6/11
675/675 [=====] - 2s 3ms/step - loss: 0.3114 -
accuracy: 0.9123
Epoch 7/11
675/675 [=====] - 2s 2ms/step - loss: 0.2968 -
accuracy: 0.9160
Epoch 8/11
675/675 [=====] - 2s 2ms/step - loss: 0.2844 -
accuracy: 0.9201
Epoch 9/11
675/675 [=====] - 2s 2ms/step - loss: 0.2738 -
accuracy: 0.9232
Epoch 10/11
675/675 [=====] - 2s 2ms/step - loss: 0.2641 -
accuracy: 0.9268
Epoch 11/11
675/675 [=====] - 2s 2ms/step - loss: 0.2555 -
accuracy: 0.9288
169/169 [=====] - 0s 2ms/step - loss: 0.2655 -
accuracy: 0.9254
Epoch 1/11
675/675 [=====] - 3s 4ms/step - loss: 1.0849 -
accuracy: 0.7234
Epoch 2/11
675/675 [=====] - 3s 4ms/step - loss: 0.4910 -
accuracy: 0.8732

```

Epoch 3/11
675/675 [=====] - 2s 3ms/step - loss: 0.4000 -
accuracy: 0.8912

Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 0.3601 -
accuracy: 0.8998

Epoch 5/11
675/675 [=====] - 2s 2ms/step - loss: 0.3354 -
accuracy: 0.9055

Epoch 6/11
675/675 [=====] - 2s 2ms/step - loss: 0.3175 -
accuracy: 0.9102

Epoch 7/11
675/675 [=====] - 2s 2ms/step - loss: 0.3032 -
accuracy: 0.9150

Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.2917 -
accuracy: 0.9183

Epoch 9/11
675/675 [=====] - 2s 2ms/step - loss: 0.2812 -
accuracy: 0.9202

Epoch 10/11
675/675 [=====] - 2s 3ms/step - loss: 0.2722 -
accuracy: 0.9234

Epoch 11/11
675/675 [=====] - 2s 3ms/step - loss: 0.2641 -
accuracy: 0.9256

169/169 [=====] - 1s 2ms/step - loss: 0.2663 -
accuracy: 0.9252

Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 1.0766 -
accuracy: 0.7332

Epoch 2/11
675/675 [=====] - 2s 2ms/step - loss: 0.4871 -
accuracy: 0.8754

Epoch 3/11
675/675 [=====] - 2s 2ms/step - loss: 0.3960 -
accuracy: 0.8933

Epoch 4/11
675/675 [=====] - 2s 2ms/step - loss: 0.3555 -
accuracy: 0.9025

Epoch 5/11
675/675 [=====] - 2s 2ms/step - loss: 0.3306 -
accuracy: 0.9082

Epoch 6/11
675/675 [=====] - 2s 2ms/step - loss: 0.3124 -
accuracy: 0.9120

Epoch 7/11

675/675 [=====] - 2s 3ms/step - loss: 0.2977 -
accuracy: 0.9155
Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.2855 -
accuracy: 0.9194
Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.2751 -
accuracy: 0.9229
Epoch 10/11
675/675 [=====] - 2s 2ms/step - loss: 0.2655 -
accuracy: 0.9253
Epoch 11/11
675/675 [=====] - 2s 2ms/step - loss: 0.2571 -
accuracy: 0.9284
169/169 [=====] - 0s 2ms/step - loss: 0.2753 -
accuracy: 0.9226
Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 1.0614 -
accuracy: 0.7369
Epoch 2/11
675/675 [=====] - 2s 2ms/step - loss: 0.4889 -
accuracy: 0.8742
Epoch 3/11
675/675 [=====] - 2s 3ms/step - loss: 0.3994 -
accuracy: 0.8912
Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 0.3586 -
accuracy: 0.9011
Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.3330 -
accuracy: 0.9073
Epoch 6/11
675/675 [=====] - 2s 2ms/step - loss: 0.3142 -
accuracy: 0.9120
Epoch 7/11
675/675 [=====] - 2s 2ms/step - loss: 0.2993 -
accuracy: 0.9159
Epoch 8/11
675/675 [=====] - 2s 2ms/step - loss: 0.2864 -
accuracy: 0.9193
Epoch 9/11
675/675 [=====] - 2s 2ms/step - loss: 0.2753 -
accuracy: 0.9224
Epoch 10/11
675/675 [=====] - 2s 2ms/step - loss: 0.2656 -
accuracy: 0.9251
Epoch 11/11
675/675 [=====] - 2s 3ms/step - loss: 0.2568 -


```

accuracy: 0.9282
169/169 [=====] - 1s 2ms/step - loss: 0.2548 -
accuracy: 0.9259
Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 0.4379 -
accuracy: 0.8767
Epoch 2/11
675/675 [=====] - 2s 2ms/step - loss: 0.2365 -
accuracy: 0.9330
Epoch 3/11
675/675 [=====] - 2s 2ms/step - loss: 0.1880 -
accuracy: 0.9458
Epoch 4/11
675/675 [=====] - 2s 2ms/step - loss: 0.1554 -
accuracy: 0.9560
Epoch 5/11
675/675 [=====] - 2s 2ms/step - loss: 0.1326 -
accuracy: 0.9619
Epoch 6/11
675/675 [=====] - 2s 2ms/step - loss: 0.1160 -
accuracy: 0.9677
Epoch 7/11
675/675 [=====] - 2s 3ms/step - loss: 0.1032 -
accuracy: 0.9705
Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.0928 -
accuracy: 0.9738
Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.0838 -
accuracy: 0.9770
Epoch 10/11
675/675 [=====] - 2s 2ms/step - loss: 0.0765 -
accuracy: 0.9782
Epoch 11/11
675/675 [=====] - 2s 2ms/step - loss: 0.0699 -
accuracy: 0.9803
169/169 [=====] - 0s 2ms/step - loss: 0.1094 -
accuracy: 0.9676
Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 0.4289 -
accuracy: 0.8800
Epoch 2/11
675/675 [=====] - 2s 2ms/step - loss: 0.2426 -
accuracy: 0.9292
Epoch 3/11
675/675 [=====] - 2s 2ms/step - loss: 0.1951 -
accuracy: 0.9438
Epoch 4/11

```

675/675 [=====] - 2s 3ms/step - loss: 0.1646 -
accuracy: 0.9533
Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.1429 -
accuracy: 0.9589
Epoch 6/11
675/675 [=====] - 2s 3ms/step - loss: 0.1266 -
accuracy: 0.9638
Epoch 7/11
675/675 [=====] - 2s 2ms/step - loss: 0.1135 -
accuracy: 0.9680
Epoch 8/11
675/675 [=====] - 2s 2ms/step - loss: 0.1039 -
accuracy: 0.9703
Epoch 9/11
675/675 [=====] - 2s 2ms/step - loss: 0.0948 -
accuracy: 0.9732
Epoch 10/11
675/675 [=====] - 2s 2ms/step - loss: 0.0871 -
accuracy: 0.9750
Epoch 11/11
675/675 [=====] - 2s 2ms/step - loss: 0.0805 -
accuracy: 0.9769
169/169 [=====] - 1s 3ms/step - loss: 0.1212 -
accuracy: 0.9652
Epoch 1/11
675/675 [=====] - 3s 3ms/step - loss: 0.4249 -
accuracy: 0.8811
Epoch 2/11
675/675 [=====] - 2s 2ms/step - loss: 0.2438 -
accuracy: 0.9299
Epoch 3/11
675/675 [=====] - 2s 2ms/step - loss: 0.1952 -
accuracy: 0.9442
Epoch 4/11
675/675 [=====] - 2s 2ms/step - loss: 0.1652 -
accuracy: 0.9524
Epoch 5/11
675/675 [=====] - 2s 2ms/step - loss: 0.1440 -
accuracy: 0.9588
Epoch 6/11
675/675 [=====] - 2s 2ms/step - loss: 0.1264 -
accuracy: 0.9636
Epoch 7/11
675/675 [=====] - 2s 2ms/step - loss: 0.1124 -
accuracy: 0.9674
Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.1012 -

```

accuracy: 0.9708
Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.0925 -
accuracy: 0.9735
Epoch 10/11
675/675 [=====] - 2s 3ms/step - loss: 0.0841 -
accuracy: 0.9756
Epoch 11/11
675/675 [=====] - 2s 2ms/step - loss: 0.0774 -
accuracy: 0.9779
169/169 [=====] - 0s 2ms/step - loss: 0.1167 -
accuracy: 0.9648
Epoch 1/11
675/675 [=====] - 2s 2ms/step - loss: 0.4281 -
accuracy: 0.8808
Epoch 2/11
675/675 [=====] - 2s 2ms/step - loss: 0.2392 -
accuracy: 0.9319
Epoch 3/11
675/675 [=====] - 2s 2ms/step - loss: 0.1911 -
accuracy: 0.9452
Epoch 4/11
675/675 [=====] - 2s 2ms/step - loss: 0.1595 -
accuracy: 0.9546
Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.1384 -
accuracy: 0.9611
Epoch 6/11
675/675 [=====] - 2s 3ms/step - loss: 0.1213 -
accuracy: 0.9656
Epoch 7/11
675/675 [=====] - 2s 3ms/step - loss: 0.1087 -
accuracy: 0.9693
Epoch 8/11
675/675 [=====] - 2s 2ms/step - loss: 0.0982 -
accuracy: 0.9720
Epoch 9/11
675/675 [=====] - 2s 2ms/step - loss: 0.0891 -
accuracy: 0.9748
Epoch 10/11
675/675 [=====] - 2s 2ms/step - loss: 0.0814 -
accuracy: 0.9775
Epoch 11/11
675/675 [=====] - 2s 2ms/step - loss: 0.0747 -
accuracy: 0.9797
169/169 [=====] - 0s 2ms/step - loss: 0.1179 -
accuracy: 0.9677
Epoch 1/11

```

675/675 [=====] - 2s 2ms/step - loss: 0.4295 -
accuracy: 0.8797
Epoch 2/11
675/675 [=====] - 2s 3ms/step - loss: 0.2461 -
accuracy: 0.9281
Epoch 3/11
675/675 [=====] - 2s 3ms/step - loss: 0.1932 -
accuracy: 0.9453
Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 0.1592 -
accuracy: 0.9547
Epoch 5/11
675/675 [=====] - 2s 2ms/step - loss: 0.1371 -
accuracy: 0.9611
Epoch 6/11
675/675 [=====] - 2s 2ms/step - loss: 0.1192 -
accuracy: 0.9666
Epoch 7/11
675/675 [=====] - 2s 2ms/step - loss: 0.1071 -
accuracy: 0.9706
Epoch 8/11
675/675 [=====] - 2s 2ms/step - loss: 0.0966 -
accuracy: 0.9733
Epoch 9/11
675/675 [=====] - 2s 2ms/step - loss: 0.0878 -
accuracy: 0.9755
Epoch 10/11
675/675 [=====] - 2s 3ms/step - loss: 0.0807 -
accuracy: 0.9780
Epoch 11/11
675/675 [=====] - 2s 3ms/step - loss: 0.0743 -
accuracy: 0.9794
169/169 [=====] - 0s 2ms/step - loss: 0.1063 -
accuracy: 0.9662
Epoch 1/11
675/675 [=====] - 2s 3ms/step - loss: 2.0628 -
accuracy: 0.3666
Epoch 2/11
675/675 [=====] - 2s 3ms/step - loss: 1.5766 -
accuracy: 0.6562
Epoch 3/11
675/675 [=====] - 3s 4ms/step - loss: 1.2372 -
accuracy: 0.7482
Epoch 4/11
675/675 [=====] - 2s 4ms/step - loss: 1.0131 -
accuracy: 0.7895
Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.8667 -

```

accuracy: 0.8114
Epoch 6/11
675/675 [=====] - 2s 3ms/step - loss: 0.7673 -
accuracy: 0.8257
Epoch 7/11
675/675 [=====] - 2s 3ms/step - loss: 0.6965 -
accuracy: 0.8365
Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.6436 -
accuracy: 0.8451
Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.6026 -
accuracy: 0.8521
Epoch 10/11
675/675 [=====] - 3s 4ms/step - loss: 0.5699 -
accuracy: 0.8578
Epoch 11/11
675/675 [=====] - 3s 4ms/step - loss: 0.5431 -
accuracy: 0.8629
169/169 [=====] - 1s 3ms/step - loss: 0.5351 -
accuracy: 0.8633
Epoch 1/11
675/675 [=====] - 2s 3ms/step - loss: 2.0653 -
accuracy: 0.3593
Epoch 2/11
675/675 [=====] - 2s 3ms/step - loss: 1.5779 -
accuracy: 0.6874
Epoch 3/11
675/675 [=====] - 2s 4ms/step - loss: 1.2184 -
accuracy: 0.7620
Epoch 4/11
675/675 [=====] - 3s 4ms/step - loss: 0.9887 -
accuracy: 0.8017
Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.8422 -
accuracy: 0.8254
Epoch 6/11
675/675 [=====] - 2s 3ms/step - loss: 0.7440 -
accuracy: 0.8388
Epoch 7/11
675/675 [=====] - 2s 3ms/step - loss: 0.6746 -
accuracy: 0.8493
Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.6231 -
accuracy: 0.8559
Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.5835 -
accuracy: 0.8615

```

Epoch 10/11
675/675 [=====] - 2s 4ms/step - loss: 0.5520 -
accuracy: 0.8663

Epoch 11/11
675/675 [=====] - 3s 4ms/step - loss: 0.5265 -
accuracy: 0.8700
169/169 [=====] - 1s 3ms/step - loss: 0.5251 -
accuracy: 0.8727

Epoch 1/11
675/675 [=====] - 2s 3ms/step - loss: 2.0326 -
accuracy: 0.3794

Epoch 2/11
675/675 [=====] - 2s 3ms/step - loss: 1.5657 -
accuracy: 0.6726

Epoch 3/11
675/675 [=====] - 2s 3ms/step - loss: 1.2235 -
accuracy: 0.7595

Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 0.9973 -
accuracy: 0.8017

Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.8510 -
accuracy: 0.8224

Epoch 6/11
675/675 [=====] - 3s 4ms/step - loss: 0.7526 -
accuracy: 0.8360

Epoch 7/11
675/675 [=====] - 3s 4ms/step - loss: 0.6829 -
accuracy: 0.8454

Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.6311 -
accuracy: 0.8524

Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.5912 -
accuracy: 0.8582

Epoch 10/11
675/675 [=====] - 2s 3ms/step - loss: 0.5595 -
accuracy: 0.8628

Epoch 11/11
675/675 [=====] - 2s 3ms/step - loss: 0.5337 -
accuracy: 0.8671
169/169 [=====] - 0s 2ms/step - loss: 0.5189 -
accuracy: 0.8702

Epoch 1/11
675/675 [=====] - 2s 3ms/step - loss: 2.0323 -
accuracy: 0.3671

Epoch 2/11
675/675 [=====] - 2s 3ms/step - loss: 1.5525 -

```

accuracy: 0.6666
Epoch 3/11
675/675 [=====] - 2s 3ms/step - loss: 1.2252 -
accuracy: 0.7535
Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 1.0076 -
accuracy: 0.7949
Epoch 5/11
675/675 [=====] - 3s 4ms/step - loss: 0.8640 -
accuracy: 0.8191
Epoch 6/11
675/675 [=====] - 3s 4ms/step - loss: 0.7655 -
accuracy: 0.8344
Epoch 7/11
675/675 [=====] - 2s 3ms/step - loss: 0.6948 -
accuracy: 0.8443
Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.6418 -
accuracy: 0.8520
Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.6007 -
accuracy: 0.8578
Epoch 10/11
675/675 [=====] - 2s 3ms/step - loss: 0.5679 -
accuracy: 0.8630
Epoch 11/11
675/675 [=====] - 2s 3ms/step - loss: 0.5412 -
accuracy: 0.8673
169/169 [=====] - 1s 3ms/step - loss: 0.5411 -
accuracy: 0.8676
Epoch 1/11
675/675 [=====] - 3s 4ms/step - loss: 2.1270 -
accuracy: 0.2780
Epoch 2/11
675/675 [=====] - 2s 3ms/step - loss: 1.6346 -
accuracy: 0.6351
Epoch 3/11
675/675 [=====] - 3s 4ms/step - loss: 1.2844 -
accuracy: 0.7383
Epoch 4/11
675/675 [=====] - 2s 4ms/step - loss: 1.0484 -
accuracy: 0.7851
Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.8919 -
accuracy: 0.8115
Epoch 6/11
675/675 [=====] - 3s 4ms/step - loss: 0.7849 -
accuracy: 0.8291

```

Epoch 7/11
675/675 [=====] - 3s 4ms/step - loss: 0.7084 -
accuracy: 0.8411

Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.6515 -
accuracy: 0.8507

Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.6077 -
accuracy: 0.8575

Epoch 10/11
675/675 [=====] - 2s 3ms/step - loss: 0.5732 -
accuracy: 0.8629

Epoch 11/11
675/675 [=====] - 2s 3ms/step - loss: 0.5450 -
accuracy: 0.8682

169/169 [=====] - 1s 3ms/step - loss: 0.5285 -
accuracy: 0.8671

Epoch 1/11
675/675 [=====] - 3s 3ms/step - loss: 1.0147 -
accuracy: 0.7592

Epoch 2/11
675/675 [=====] - 2s 3ms/step - loss: 0.4802 -
accuracy: 0.8750

Epoch 3/11
675/675 [=====] - 2s 3ms/step - loss: 0.3940 -
accuracy: 0.8924

Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 0.3536 -
accuracy: 0.9016

Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.3278 -
accuracy: 0.9083

Epoch 6/11
675/675 [=====] - 3s 4ms/step - loss: 0.3089 -
accuracy: 0.9134

Epoch 7/11
675/675 [=====] - 3s 4ms/step - loss: 0.2935 -
accuracy: 0.9178

Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.2801 -
accuracy: 0.9218

Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.2689 -
accuracy: 0.9253

Epoch 10/11
675/675 [=====] - 2s 3ms/step - loss: 0.2587 -
accuracy: 0.9279

Epoch 11/11


```

675/675 [=====] - 2s 3ms/step - loss: 0.2491 -
accuracy: 0.9305
169/169 [=====] - 0s 2ms/step - loss: 0.2563 -
accuracy: 0.9289
Epoch 1/11
675/675 [=====] - 3s 3ms/step - loss: 1.0078 -
accuracy: 0.7593
Epoch 2/11
675/675 [=====] - 3s 4ms/step - loss: 0.4689 -
accuracy: 0.8787
Epoch 3/11
675/675 [=====] - 3s 4ms/step - loss: 0.3856 -
accuracy: 0.8945
Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 0.3467 -
accuracy: 0.9041
Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.3217 -
accuracy: 0.9108
Epoch 6/11
675/675 [=====] - 2s 3ms/step - loss: 0.3031 -
accuracy: 0.9153
Epoch 7/11
675/675 [=====] - 2s 3ms/step - loss: 0.2879 -
accuracy: 0.9191
Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.2750 -
accuracy: 0.9234
Epoch 9/11
675/675 [=====] - 3s 4ms/step - loss: 0.2643 -
accuracy: 0.9263
Epoch 10/11
675/675 [=====] - 3s 4ms/step - loss: 0.2543 -
accuracy: 0.9293
Epoch 11/11
675/675 [=====] - 2s 3ms/step - loss: 0.2453 -
accuracy: 0.9317
169/169 [=====] - 0s 2ms/step - loss: 0.2521 -
accuracy: 0.9301
Epoch 1/11
675/675 [=====] - 3s 3ms/step - loss: 1.0027 -
accuracy: 0.7581
Epoch 2/11
675/675 [=====] - 2s 3ms/step - loss: 0.4702 -
accuracy: 0.8791
Epoch 3/11
675/675 [=====] - 2s 3ms/step - loss: 0.3891 -
accuracy: 0.8943

```

Epoch 4/11
675/675 [=====] - 3s 4ms/step - loss: 0.3516 -
accuracy: 0.9033

Epoch 5/11
675/675 [=====] - 3s 4ms/step - loss: 0.3274 -
accuracy: 0.9099

Epoch 6/11
675/675 [=====] - 2s 3ms/step - loss: 0.3095 -
accuracy: 0.9137

Epoch 7/11
675/675 [=====] - 2s 3ms/step - loss: 0.2947 -
accuracy: 0.9179

Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.2824 -
accuracy: 0.9218

Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.2713 -
accuracy: 0.9246

Epoch 10/11
675/675 [=====] - 2s 3ms/step - loss: 0.2614 -
accuracy: 0.9279

Epoch 11/11
675/675 [=====] - 3s 4ms/step - loss: 0.2523 -
accuracy: 0.9293

169/169 [=====] - 1s 3ms/step - loss: 0.2554 -
accuracy: 0.9271

Epoch 1/11
675/675 [=====] - 3s 3ms/step - loss: 1.0089 -
accuracy: 0.7618

Epoch 2/11
675/675 [=====] - 2s 3ms/step - loss: 0.4662 -
accuracy: 0.8786

Epoch 3/11
675/675 [=====] - 2s 3ms/step - loss: 0.3840 -
accuracy: 0.8946

Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 0.3448 -
accuracy: 0.9039

Epoch 5/11
675/675 [=====] - 2s 4ms/step - loss: 0.3195 -
accuracy: 0.9104

Epoch 6/11
675/675 [=====] - 3s 4ms/step - loss: 0.3009 -
accuracy: 0.9147

Epoch 7/11
675/675 [=====] - 3s 4ms/step - loss: 0.2857 -
accuracy: 0.9194

Epoch 8/11

```

675/675 [=====] - 2s 3ms/step - loss: 0.2733 -
accuracy: 0.9233
Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.2622 -
accuracy: 0.9263
Epoch 10/11
675/675 [=====] - 2s 3ms/step - loss: 0.2526 -
accuracy: 0.9286
Epoch 11/11
675/675 [=====] - 2s 3ms/step - loss: 0.2440 -
accuracy: 0.9314
169/169 [=====] - 0s 2ms/step - loss: 0.2637 -
accuracy: 0.9250
Epoch 1/11
675/675 [=====] - 3s 4ms/step - loss: 0.9979 -
accuracy: 0.7600
Epoch 2/11
675/675 [=====] - 3s 4ms/step - loss: 0.4617 -
accuracy: 0.8806
Epoch 3/11
675/675 [=====] - 2s 3ms/step - loss: 0.3845 -
accuracy: 0.8956
Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 0.3486 -
accuracy: 0.9034
Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.3258 -
accuracy: 0.9088
Epoch 6/11
675/675 [=====] - 2s 3ms/step - loss: 0.3088 -
accuracy: 0.9136
Epoch 7/11
675/675 [=====] - 2s 3ms/step - loss: 0.2953 -
accuracy: 0.9166
Epoch 8/11
675/675 [=====] - 3s 4ms/step - loss: 0.2839 -
accuracy: 0.9198
Epoch 9/11
675/675 [=====] - 3s 4ms/step - loss: 0.2737 -
accuracy: 0.9227
Epoch 10/11
675/675 [=====] - 2s 3ms/step - loss: 0.2644 -
accuracy: 0.9250
Epoch 11/11
675/675 [=====] - 3s 4ms/step - loss: 0.2561 -
accuracy: 0.9281
169/169 [=====] - 0s 2ms/step - loss: 0.2556 -
accuracy: 0.9273

```

Epoch 1/11
675/675 [=====] - 3s 4ms/step - loss: 0.4156 -
accuracy: 0.8858

Epoch 2/11
675/675 [=====] - 3s 5ms/step - loss: 0.2278 -
accuracy: 0.9348

Epoch 3/11
675/675 [=====] - 2s 3ms/step - loss: 0.1756 -
accuracy: 0.9497

Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 0.1434 -
accuracy: 0.9601

Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.1214 -
accuracy: 0.9652

Epoch 6/11
675/675 [=====] - 2s 3ms/step - loss: 0.1052 -
accuracy: 0.9700

Epoch 7/11
675/675 [=====] - 4s 6ms/step - loss: 0.0921 -
accuracy: 0.9738

Epoch 8/11
675/675 [=====] - 3s 5ms/step - loss: 0.0819 -
accuracy: 0.9770

Epoch 9/11
675/675 [=====] - 3s 4ms/step - loss: 0.0737 -
accuracy: 0.9796

Epoch 10/11
675/675 [=====] - 2s 3ms/step - loss: 0.0664 -
accuracy: 0.9817

Epoch 11/11
675/675 [=====] - 2s 3ms/step - loss: 0.0600 -
accuracy: 0.9838

169/169 [=====] - 0s 2ms/step - loss: 0.1043 -
accuracy: 0.9705

Epoch 1/11
675/675 [=====] - 2s 3ms/step - loss: 0.4232 -
accuracy: 0.8829

Epoch 2/11
675/675 [=====] - 2s 3ms/step - loss: 0.2268 -
accuracy: 0.9357

Epoch 3/11
675/675 [=====] - 3s 4ms/step - loss: 0.1750 -
accuracy: 0.9504

Epoch 4/11
675/675 [=====] - 3s 4ms/step - loss: 0.1438 -
accuracy: 0.9590

Epoch 5/11

675/675 [=====] - 2s 3ms/step - loss: 0.1225 -
accuracy: 0.9655
Epoch 6/11
675/675 [=====] - 2s 3ms/step - loss: 0.1059 -
accuracy: 0.9704
Epoch 7/11
675/675 [=====] - 2s 3ms/step - loss: 0.0939 -
accuracy: 0.9740
Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.0837 -
accuracy: 0.9765
Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.0748 -
accuracy: 0.9797
Epoch 10/11
675/675 [=====] - 3s 4ms/step - loss: 0.0678 -
accuracy: 0.9812
Epoch 11/11
675/675 [=====] - 3s 4ms/step - loss: 0.0616 -
accuracy: 0.9823
169/169 [=====] - 0s 2ms/step - loss: 0.0978 -
accuracy: 0.9719
Epoch 1/11
675/675 [=====] - 2s 3ms/step - loss: 0.4208 -
accuracy: 0.8840
Epoch 2/11
675/675 [=====] - 2s 3ms/step - loss: 0.2305 -
accuracy: 0.9344
Epoch 3/11
675/675 [=====] - 2s 3ms/step - loss: 0.1785 -
accuracy: 0.9500
Epoch 4/11
675/675 [=====] - 3s 4ms/step - loss: 0.1461 -
accuracy: 0.9574
Epoch 5/11
675/675 [=====] - 2s 4ms/step - loss: 0.1227 -
accuracy: 0.9660
Epoch 6/11
675/675 [=====] - 2s 3ms/step - loss: 0.1065 -
accuracy: 0.9700
Epoch 7/11
675/675 [=====] - 2s 3ms/step - loss: 0.0933 -
accuracy: 0.9734
Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.0829 -
accuracy: 0.9773
Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.0743 -

```

accuracy: 0.9798
Epoch 10/11
675/675 [=====] - 2s 3ms/step - loss: 0.0668 -
accuracy: 0.9817
Epoch 11/11
675/675 [=====] - 3s 4ms/step - loss: 0.0607 -
accuracy: 0.9840
169/169 [=====] - 1s 3ms/step - loss: 0.0947 -
accuracy: 0.9735
Epoch 1/11
675/675 [=====] - 2s 3ms/step - loss: 0.4196 -
accuracy: 0.8834
Epoch 2/11
675/675 [=====] - 2s 3ms/step - loss: 0.2278 -
accuracy: 0.9352
Epoch 3/11
675/675 [=====] - 2s 3ms/step - loss: 0.1742 -
accuracy: 0.9506
Epoch 4/11
675/675 [=====] - 2s 3ms/step - loss: 0.1421 -
accuracy: 0.9592
Epoch 5/11
675/675 [=====] - 2s 3ms/step - loss: 0.1202 -
accuracy: 0.9666
Epoch 6/11
675/675 [=====] - 3s 4ms/step - loss: 0.1043 -
accuracy: 0.9712
Epoch 7/11
675/675 [=====] - 3s 4ms/step - loss: 0.0916 -
accuracy: 0.9740
Epoch 8/11
675/675 [=====] - 2s 3ms/step - loss: 0.0815 -
accuracy: 0.9773
Epoch 9/11
675/675 [=====] - 2s 3ms/step - loss: 0.0733 -
accuracy: 0.9795
Epoch 10/11
675/675 [=====] - 2s 3ms/step - loss: 0.0664 -
accuracy: 0.9822
Epoch 11/11
675/675 [=====] - 2s 3ms/step - loss: 0.0603 -
accuracy: 0.9840
169/169 [=====] - 0s 2ms/step - loss: 0.1022 -
accuracy: 0.9686
Epoch 1/11
675/675 [=====] - 2s 3ms/step - loss: 0.4147 -
accuracy: 0.8861
Epoch 2/11

```

675/675 [=====] - 2s 4ms/step - loss: 0.2292 -
 accuracy: 0.9354
 Epoch 3/11
 675/675 [=====] - 3s 4ms/step - loss: 0.1780 -
 accuracy: 0.9493
 Epoch 4/11
 675/675 [=====] - 3s 4ms/step - loss: 0.1448 -
 accuracy: 0.9589
 Epoch 5/11
 675/675 [=====] - 3s 4ms/step - loss: 0.1235 -
 accuracy: 0.9655
 Epoch 6/11
 675/675 [=====] - 3s 4ms/step - loss: 0.1069 -
 accuracy: 0.9705
 Epoch 7/11
 675/675 [=====] - 2s 3ms/step - loss: 0.0933 -
 accuracy: 0.9744
 Epoch 8/11
 675/675 [=====] - 2s 3ms/step - loss: 0.0836 -
 accuracy: 0.9766
 Epoch 9/11
 675/675 [=====] - 2s 3ms/step - loss: 0.0740 -
 accuracy: 0.9801
 Epoch 10/11
 675/675 [=====] - 2s 3ms/step - loss: 0.0678 -
 accuracy: 0.9811
 Epoch 11/11
 675/675 [=====] - 2s 3ms/step - loss: 0.0614 -
 accuracy: 0.9834
 169/169 [=====] - 1s 3ms/step - loss: 0.0923 -
 accuracy: 0.9711
 Epoch 1/11
 844/844 [=====] - 4s 4ms/step - loss: 0.3869 -
 accuracy: 0.8926
 Epoch 2/11
 844/844 [=====] - 3s 3ms/step - loss: 0.2067 -
 accuracy: 0.9414
 Epoch 3/11
 844/844 [=====] - 3s 3ms/step - loss: 0.1578 -
 accuracy: 0.9552
 Epoch 4/11
 844/844 [=====] - 3s 3ms/step - loss: 0.1283 -
 accuracy: 0.9639
 Epoch 5/11
 844/844 [=====] - 3s 4ms/step - loss: 0.1091 -
 accuracy: 0.9694
 Epoch 6/11
 844/844 [=====] - 4s 4ms/step - loss: 0.0947 -

```

accuracy: 0.9735
Epoch 7/11
844/844 [=====] - 3s 3ms/step - loss: 0.0833 -
accuracy: 0.9770
Epoch 8/11
844/844 [=====] - 3s 3ms/step - loss: 0.0736 -
accuracy: 0.9797
Epoch 9/11
844/844 [=====] - 3s 3ms/step - loss: 0.0666 -
accuracy: 0.9815
Epoch 10/11
844/844 [=====] - 2s 3ms/step - loss: 0.0604 -
accuracy: 0.9830
Epoch 11/11
844/844 [=====] - 3s 4ms/step - loss: 0.0545 -
accuracy: 0.9852

```

```

Best Hyperparameters: {'hidden_units': 128, 'learning_rate': 0.1}
Best Accuracy: 0.9711111068725586

```

```


[ ]: # Get the best model from the grid search
best_model = grid_result.best_estimator_

# Get the training history
history = best_model.fit(x_train, y_train, epochs=11, batch_size=64, verbose=0,
    ↪validation_data=(x_val, y_val))

# Plot the learning curve
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Learning Curve')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

print("\n\n")
# Print the best hyperparameters and accuracy
print("Best Hyperparameters: ", grid_result.best_params_)
print("Best Accuracy: ", grid_result.best_score_)

```

Best Hyperparameters: `{'hidden_units': 128, 'learning_rate': 0.1}`
Best Accuracy: 0.9711111068725586

In this experiment, our objective was to construct a simple neural network comprising only one hidden layer. To optimize its performance, we employed GridSearchCV to explore various values for the learning rate and the number of neurons in the hidden layer. In addition, we selected 'accuracy' as the evaluation metric and utilized cross-entropy categorical loss as our loss function, which is suitable for multiclass classification scenarios.

To facilitate learning, we employed mini-batch stochastic gradient descent (SGD) as our chosen learning algorithm, employing a batch size of 64. Furthermore, we integrated a validation set to conduct ongoing validation during the training process. In the results, we can see the best accuracy achieved is 97.1% with hidden units as 128 and learning rate of 0.1.

1.2 Que 1: Part 2

-
2. It is recommended to try, at least, another optimization method of your choice (SGD with

momentum, RMSProp, RMSProp with momentum, AdaGrad, AdaDelta, or Adam) and compare its performances to those of the basic minibatch SGD on the MNIST dataset. Which methods you want to try and how many you want to try and compare is up to you and up to the amount of time you have left to complete the assignment. Remember, this is a research course. You may want to read Chapter-8 also.

```
[ ]: import matplotlib.pyplot as plt
from tensorflow.keras.optimizers import SGD, RMSprop, Adagrad, Adadelta, Adam

learning_rates = [0.001, 0.01, 0.1]
optimizers = [
    ('SGD', SGD),
    ('RMSprop', RMSprop),
    ('Adagrad', Adagrad),
    ('Adadelta', Adadelta),
    ('Adam', Adam)
]
# Find the best learning rate and optimizer based on model accuracy
best_accuracy = 0.0
best_learning_rate = None
best_optimizer = None

# Define the model architecture
model = Sequential()
model.add(Dense(128, activation='relu', input_shape=(784,)))
model.add(Dense(10, activation='sigmoid'))

# Train and plot accuracy for different learning rates and optimizers
fig, axs = plt.subplots(3, len(optimizers), figsize=(12, 12))

for i, lr in enumerate(learning_rates):
    for j, (optimizer_name, optimizer_class) in enumerate(optimizers):
        optimizer = optimizer_class(learning_rate=lr)
        model.compile(loss='categorical_crossentropy', optimizer=optimizer,
            ↪metrics=['accuracy'])
        history = model.fit(x_train, y_train, batch_size=64, epochs=11,
            ↪validation_split=0.2, verbose=0)
        accuracy = history.history['val_accuracy'][-1]

        if accuracy > best_accuracy:
            best_accuracy = accuracy
            best_learning_rate = lr
            best_optimizer = optimizer_name

    # Plot accuracy
    axs[i, j].plot(history.history['val_accuracy'])
```

```

    # Set plot labels
    axs[i, j].set_title(f'Optimizer: {optimizer_name}\nLearning Rate: {lr}')
    axs[i, j].set_xlabel('Epochs')
    axs[i, j].set_ylabel('Validation Accuracy')

# Set layout and spacing between subplots
plt.tight_layout(pad=2.0)

# Print the best learning rate and optimizer
print(f'Best Learning Rate: {best_learning_rate}')
print(f'Best Optimizer: {best_optimizer}')
print(f'Best Accuracy: {best_accuracy}')

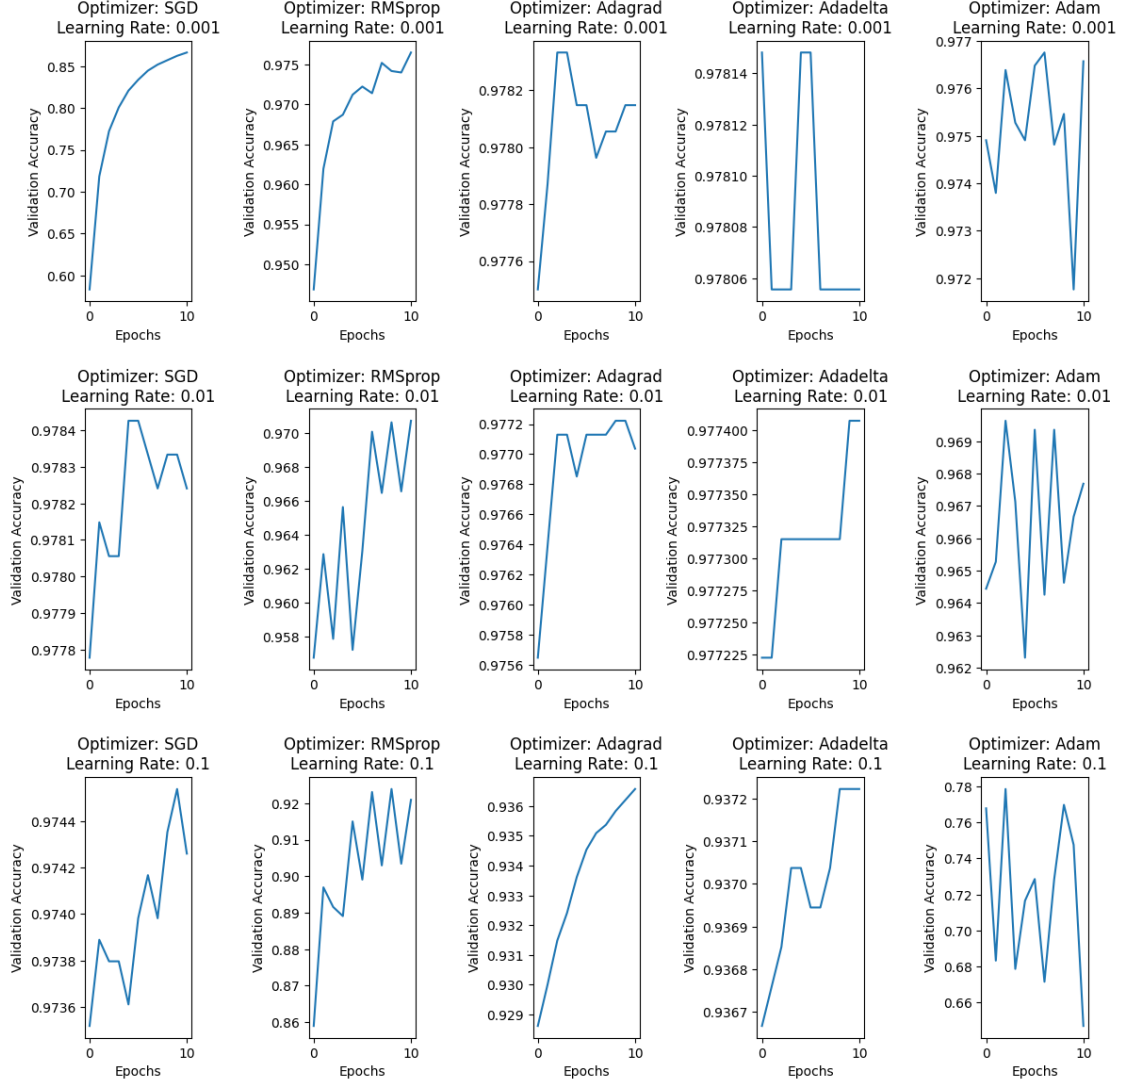
print("\n\n")
# Show the plots
plt.show()

```

```

Best Learning Rate: 0.01
Best Optimizer: SGD
Best Accuracy: 0.9782407283782959

```



In our experiment, we investigated the performance of various optimization algorithms by testing them with three different learning rate values: 0.1, 0.01, and 0.001. To ensure a fair comparison, we kept the batch size, number of epochs, loss function, number of layers, and hidden units consistent across all methods.

By analyzing the plots of validation accuracy versus epochs for each optimization algorithm, we can draw some conclusions.

Firstly, both SGD and RMSProp consistently outperformed Adagrad and Adadelta. This observation suggests that SGD and RMSProp were more effective in finding the optimal parameters for our neural network. The reason behind their superior performance could be attributed to their ability to adapt the learning rate based on the gradients of the parameters.

On the other hand, the performance of Adam exhibited high variation in the plots. This variation could be attributed to a couple of factors. One possibility is that we trained the neural network for

a relatively small number of epochs, which might not have been sufficient for Adam to converge to a stable solution. Additionally, Adam incorporates adaptive learning rates and momentum, which can introduce additional complexity in optimization, especially in scenarios with simple data.

Despite the variations observed with Adam, it is worth noting that all the optimization algorithms achieved an accuracy of over 90 percent. This outcome suggests that the dataset used in our experiment was relatively simple, allowing all the algorithms to achieve satisfactory results. It is important to consider that more complex datasets may require further fine-tuning of hyperparameters and choice of optimization algorithm to achieve optimal performance.

Overall, through these iterations we achieved 97.8% as a best accuracy and optimizer was “SGD” with learning rate 0.01.

2 Que 2:

Consider the L2-regularized multiclass logistic regression. That is, add to the logistic regression loss a regularization term that represents the L2-norm of the parameters. More precisely, the regularization term is

$$(w, b) = \lambda (w^2 + b^2)$$

where $\{w_i, b_i\}$ are all the parameters in the logistic regression, and λ is the regularization hyperparameter. Typically, λ is about C/n where n is the number of data points and C is some constant in $[0.01, 100]$ (need to tune C). Run the regularized multiclass logistic regression on MNIST, using the basic minibatch SGD, and compare its results to those of the basic minibatch SGD with nonregularized loss, in Question #1.

```
[25]: import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, InputLayer
from tensorflow.keras.optimizers import SGD
from sklearn.preprocessing import OneHotEncoder

# Define the logistic regression model architecture with L2 regularization
model = Sequential()
model.add(Dense(128, activation='relu', input_shape=(784,),
    ↪kernel_regularizer='l2'))
model.add(Dense(10, activation='sigmoid'))

# Compile the model
model.compile(loss='categorical_crossentropy', optimizer=SGD(learning_rate=0.
    ↪1), metrics=['accuracy'])

# Train the model with L2 regularization
history = model.fit(x_train, y_train, batch_size=128, epochs=11,
    ↪validation_split=0.2, verbose=1)
```

```

# Evaluate the model on the testing set
_, test_accuracy = model.evaluate(x_test, y_test, verbose=0)

print("\n")
# Plot the training and validation accuracies
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('L2-Regularized Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(['Train', 'Validation'])
plt.show()

print("\n\n")
# Print the accuracy
print(f'Test Accuracy (L2-Regularized): {test_accuracy}')

```


```

Epoch 1/11
338/338 [=====] - 2s 5ms/step - loss: 1.8499 -
accuracy: 0.8566 - val_loss: 1.0354 - val_accuracy: 0.8985
Epoch 2/11
338/338 [=====] - 2s 4ms/step - loss: 0.7551 -
accuracy: 0.9112 - val_loss: 0.5689 - val_accuracy: 0.9146
Epoch 3/11
338/338 [=====] - 2s 5ms/step - loss: 0.4855 -
accuracy: 0.9212 - val_loss: 0.4373 - val_accuracy: 0.9202
Epoch 4/11
338/338 [=====] - 2s 6ms/step - loss: 0.4004 -
accuracy: 0.9283 - val_loss: 0.3874 - val_accuracy: 0.9295
Epoch 5/11
338/338 [=====] - 2s 7ms/step - loss: 0.3652 -
accuracy: 0.9319 - val_loss: 0.3551 - val_accuracy: 0.9307
Epoch 6/11
338/338 [=====] - 2s 5ms/step - loss: 0.3454 -
accuracy: 0.9350 - val_loss: 0.3500 - val_accuracy: 0.9322
Epoch 7/11
338/338 [=====] - 1s 4ms/step - loss: 0.3320 -
accuracy: 0.9388 - val_loss: 0.3371 - val_accuracy: 0.9372
Epoch 8/11
338/338 [=====] - 1s 4ms/step - loss: 0.3207 -
accuracy: 0.9415 - val_loss: 0.3377 - val_accuracy: 0.9331
Epoch 9/11
338/338 [=====] - 1s 4ms/step - loss: 0.3136 -
accuracy: 0.9418 - val_loss: 0.3129 - val_accuracy: 0.9409
Epoch 10/11
338/338 [=====] - 1s 4ms/step - loss: 0.3057 -
accuracy: 0.9441 - val_loss: 0.3383 - val_accuracy: 0.9328

```

Epoch 11/11

338/338 [=====] - 2s 4ms/step - loss: 0.2990 -
accuracy: 0.9461 - val_loss: 0.3043 - val_accuracy: 0.9439



Test Accuracy (L2-Regularized): 0.9488000273704529

We applied L2 regularization to the same model that was used in the previous question. L2 regularization is a technique commonly employed to prevent overfitting in neural networks by adding a penalty term to the loss function based on the magnitudes of the model's weights.

After incorporating L2 regularization, we observed a slight decrease in the training accuracy compared to the previous result. The training accuracy now stands at 94.62% instead of the previous 97.1%. This drop in training accuracy can be attributed to the regularization term penalizing the weights and encouraging them to be smaller. While we achieved 94.88% as a test accuracy with L2-Regularization.

However, an interesting observation is that the validation accuracy is now closer to the training

accuracy after the addition of L2 regularization. This outcome suggests that the regularization helped in reducing overfitting and improving the generalization capability of the model. When the validation accuracy is closer to the training accuracy, it indicates that the model's performance on unseen data is more aligned with its performance on the training data.

Overall, although the training accuracy decreased slightly due to the L2 regularization, the improved alignment between the training and validation accuracies implies that the model is now less likely to overfit and more capable of generalizing well to unseen data.

3 Que 3:

Build a three-layer feedforward neural network:

$x \rightarrow h1 \rightarrow h2 \rightarrow o$

The hidden layers h1 and h2 have width 500. Train the network for 250 epochs¹ and test the classification error. Do not use regularizations. Plot the cross-entropy loss on the batches and also plot the classification error on the validation data.

Comments: ¹Each epoch is a pass over the training data. Suppose you use batches of size b, and the training data set has n points, then an epoch consists of n/b batches. Note that you can divide the data set into batches, and then round robin over the batches. You can also randomly sample say 64 points for each batch. Either way is OK, and typically there is no performance difference between them. When these batches are randomly sampled, it is possible that some point are not in any of them, but we still call these batches a pass over the data.

Comments: you can also use another dataset, CIFAR-10 or CIFAR-100. Or you can pick your own dataset.

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import SGD
from sklearn.preprocessing import OneHotEncoder
from sklearn.metrics import accuracy_score

# Define the three-layer feedforward neural network
model = Sequential()
model.add(Dense(500, activation='relu', input_shape=(784,)))
model.add(Dense(500, activation='relu'))
model.add(Dense(10, activation='sigmoid'))

# Compile the model
model.compile(loss='categorical_crossentropy', optimizer=SGD(learning_rate=0.
    ↪01), metrics=['accuracy'])

# Train the model
```



```

history = model.fit(x_train, y_train, batch_size=64, epochs=250,
                    validation_split=0.2, verbose=1)

# Evaluate the model on the testing set
_, test_accuracy = model.evaluate(x_test, y_test, verbose=0)

print("\n")
# Plot the cross-entropy loss on the batches
plt.plot(history.history['loss'])
plt.title('Cross-Entropy Loss on Batches')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.show()
print("\n")

# Plot the classification error on the validation data
plt.plot(1 - np.array(history.history['val_accuracy']))
plt.title('Classification Error on Validation Data')
plt.xlabel('Epoch')
plt.ylabel('Error')
plt.show()

print("\n\n")
# Print the test accuracy
print(f'Test Accuracy: {test_accuracy}')

```

```

Epoch 1/250
675/675 [=====] - 7s 9ms/step - loss: 0.9166 -
accuracy: 0.7875 - val_loss: 0.4505 - val_accuracy: 0.8779
Epoch 2/250
675/675 [=====] - 8s 11ms/step - loss: 0.3909 -
accuracy: 0.8938 - val_loss: 0.3414 - val_accuracy: 0.9004
Epoch 3/250
675/675 [=====] - 6s 9ms/step - loss: 0.3231 -
accuracy: 0.9090 - val_loss: 0.3014 - val_accuracy: 0.9111
Epoch 4/250
675/675 [=====] - 7s 11ms/step - loss: 0.2884 -
accuracy: 0.9179 - val_loss: 0.2765 - val_accuracy: 0.9182
Epoch 5/250
675/675 [=====] - 7s 10ms/step - loss: 0.2635 -
accuracy: 0.9257 - val_loss: 0.2545 - val_accuracy: 0.9285
Epoch 6/250
675/675 [=====] - 6s 9ms/step - loss: 0.2439 -
accuracy: 0.9312 - val_loss: 0.2393 - val_accuracy: 0.9319
Epoch 7/250
675/675 [=====] - 9s 14ms/step - loss: 0.2270 -
accuracy: 0.9358 - val_loss: 0.2276 - val_accuracy: 0.9352

```

Epoch 8/250
675/675 [=====] - 7s 10ms/step - loss: 0.2129 -
accuracy: 0.9401 - val_loss: 0.2151 - val_accuracy: 0.9385
Epoch 9/250
675/675 [=====] - 7s 10ms/step - loss: 0.2002 -
accuracy: 0.9439 - val_loss: 0.2035 - val_accuracy: 0.9415
Epoch 10/250
675/675 [=====] - 7s 10ms/step - loss: 0.1888 -
accuracy: 0.9471 - val_loss: 0.1958 - val_accuracy: 0.9421
Epoch 11/250
675/675 [=====] - 7s 10ms/step - loss: 0.1787 -
accuracy: 0.9497 - val_loss: 0.1855 - val_accuracy: 0.9452
Epoch 12/250
675/675 [=====] - 6s 9ms/step - loss: 0.1695 -
accuracy: 0.9525 - val_loss: 0.1780 - val_accuracy: 0.9487
Epoch 13/250
675/675 [=====] - 8s 12ms/step - loss: 0.1606 -
accuracy: 0.9555 - val_loss: 0.1710 - val_accuracy: 0.9503
Epoch 14/250
675/675 [=====] - 6s 9ms/step - loss: 0.1533 -
accuracy: 0.9573 - val_loss: 0.1660 - val_accuracy: 0.9497
Epoch 15/250
675/675 [=====] - 8s 11ms/step - loss: 0.1459 -
accuracy: 0.9594 - val_loss: 0.1581 - val_accuracy: 0.9542
Epoch 16/250
675/675 [=====] - 6s 9ms/step - loss: 0.1392 -
accuracy: 0.9614 - val_loss: 0.1555 - val_accuracy: 0.9564
Epoch 17/250
675/675 [=====] - 8s 12ms/step - loss: 0.1330 -
accuracy: 0.9637 - val_loss: 0.1500 - val_accuracy: 0.9559
Epoch 18/250
675/675 [=====] - 6s 9ms/step - loss: 0.1273 -
accuracy: 0.9651 - val_loss: 0.1464 - val_accuracy: 0.9574
Epoch 19/250
675/675 [=====] - 8s 12ms/step - loss: 0.1220 -
accuracy: 0.9670 - val_loss: 0.1400 - val_accuracy: 0.9597
Epoch 20/250
675/675 [=====] - 6s 9ms/step - loss: 0.1169 -
accuracy: 0.9685 - val_loss: 0.1373 - val_accuracy: 0.9593
Epoch 21/250
675/675 [=====] - 8s 12ms/step - loss: 0.1120 -
accuracy: 0.9699 - val_loss: 0.1345 - val_accuracy: 0.9604
Epoch 22/250
675/675 [=====] - 6s 9ms/step - loss: 0.1078 -
accuracy: 0.9714 - val_loss: 0.1316 - val_accuracy: 0.9608
Epoch 23/250
675/675 [=====] - 8s 11ms/step - loss: 0.1034 -
accuracy: 0.9727 - val_loss: 0.1289 - val_accuracy: 0.9619

Epoch 24/250
675/675 [=====] - 6s 9ms/step - loss: 0.0998 - accuracy: 0.9732 - val_loss: 0.1242 - val_accuracy: 0.9632

Epoch 25/250
675/675 [=====] - 7s 11ms/step - loss: 0.0960 - accuracy: 0.9746 - val_loss: 0.1221 - val_accuracy: 0.9637

Epoch 26/250
675/675 [=====] - 7s 10ms/step - loss: 0.0922 - accuracy: 0.9755 - val_loss: 0.1218 - val_accuracy: 0.9637

Epoch 27/250
675/675 [=====] - 7s 10ms/step - loss: 0.0890 - accuracy: 0.9764 - val_loss: 0.1190 - val_accuracy: 0.9639

Epoch 28/250
675/675 [=====] - 7s 10ms/step - loss: 0.0856 - accuracy: 0.9778 - val_loss: 0.1155 - val_accuracy: 0.9656

Epoch 29/250
675/675 [=====] - 6s 9ms/step - loss: 0.0828 - accuracy: 0.9780 - val_loss: 0.1146 - val_accuracy: 0.9641

Epoch 30/250
675/675 [=====] - 10s 14ms/step - loss: 0.0799 - accuracy: 0.9789 - val_loss: 0.1105 - val_accuracy: 0.9671

Epoch 31/250
675/675 [=====] - 7s 11ms/step - loss: 0.0772 - accuracy: 0.9799 - val_loss: 0.1097 - val_accuracy: 0.9670

Epoch 32/250
675/675 [=====] - 7s 10ms/step - loss: 0.0745 - accuracy: 0.9803 - val_loss: 0.1094 - val_accuracy: 0.9667

Epoch 33/250
675/675 [=====] - 7s 10ms/step - loss: 0.0718 - accuracy: 0.9814 - val_loss: 0.1078 - val_accuracy: 0.9683

Epoch 34/250
675/675 [=====] - 7s 10ms/step - loss: 0.0695 - accuracy: 0.9822 - val_loss: 0.1039 - val_accuracy: 0.9692

Epoch 35/250
675/675 [=====] - 6s 9ms/step - loss: 0.0673 - accuracy: 0.9823 - val_loss: 0.1047 - val_accuracy: 0.9679

Epoch 36/250
675/675 [=====] - 8s 11ms/step - loss: 0.0651 - accuracy: 0.9837 - val_loss: 0.1015 - val_accuracy: 0.9700

Epoch 37/250
675/675 [=====] - 6s 9ms/step - loss: 0.0629 - accuracy: 0.9847 - val_loss: 0.1006 - val_accuracy: 0.9695

Epoch 38/250
675/675 [=====] - 8s 12ms/step - loss: 0.0609 - accuracy: 0.9848 - val_loss: 0.1007 - val_accuracy: 0.9694

Epoch 39/250
675/675 [=====] - 6s 9ms/step - loss: 0.0588 - accuracy: 0.9856 - val_loss: 0.0997 - val_accuracy: 0.9698

Epoch 40/250
675/675 [=====] - 8s 12ms/step - loss: 0.0570 -
accuracy: 0.9860 - val_loss: 0.0980 - val_accuracy: 0.9700
Epoch 41/250
675/675 [=====] - 6s 9ms/step - loss: 0.0553 -
accuracy: 0.9862 - val_loss: 0.0971 - val_accuracy: 0.9710
Epoch 42/250
675/675 [=====] - 8s 12ms/step - loss: 0.0535 -
accuracy: 0.9870 - val_loss: 0.0955 - val_accuracy: 0.9706
Epoch 43/250
675/675 [=====] - 6s 9ms/step - loss: 0.0518 -
accuracy: 0.9876 - val_loss: 0.0956 - val_accuracy: 0.9702
Epoch 44/250
675/675 [=====] - 8s 12ms/step - loss: 0.0501 -
accuracy: 0.9881 - val_loss: 0.0952 - val_accuracy: 0.9702
Epoch 45/250
675/675 [=====] - 6s 9ms/step - loss: 0.0486 -
accuracy: 0.9883 - val_loss: 0.0933 - val_accuracy: 0.9708
Epoch 46/250
675/675 [=====] - 8s 11ms/step - loss: 0.0472 -
accuracy: 0.9891 - val_loss: 0.0930 - val_accuracy: 0.9711
Epoch 47/250
675/675 [=====] - 6s 10ms/step - loss: 0.0456 -
accuracy: 0.9891 - val_loss: 0.0926 - val_accuracy: 0.9709
Epoch 48/250
675/675 [=====] - 7s 11ms/step - loss: 0.0444 -
accuracy: 0.9900 - val_loss: 0.0916 - val_accuracy: 0.9702
Epoch 49/250
675/675 [=====] - 6s 10ms/step - loss: 0.0431 -
accuracy: 0.9901 - val_loss: 0.0909 - val_accuracy: 0.9714
Epoch 50/250
675/675 [=====] - 7s 10ms/step - loss: 0.0418 -
accuracy: 0.9907 - val_loss: 0.0898 - val_accuracy: 0.9720
Epoch 51/250
675/675 [=====] - 7s 10ms/step - loss: 0.0406 -
accuracy: 0.9907 - val_loss: 0.0899 - val_accuracy: 0.9726
Epoch 52/250
675/675 [=====] - 6s 9ms/step - loss: 0.0394 -
accuracy: 0.9911 - val_loss: 0.0901 - val_accuracy: 0.9706
Epoch 53/250
675/675 [=====] - 8s 11ms/step - loss: 0.0383 -
accuracy: 0.9918 - val_loss: 0.0890 - val_accuracy: 0.9716
Epoch 54/250
675/675 [=====] - 9s 13ms/step - loss: 0.0371 -
accuracy: 0.9918 - val_loss: 0.0890 - val_accuracy: 0.9719
Epoch 55/250
675/675 [=====] - 7s 11ms/step - loss: 0.0361 -
accuracy: 0.9923 - val_loss: 0.0882 - val_accuracy: 0.9722

Epoch 56/250
675/675 [=====] - 7s 10ms/step - loss: 0.0349 - accuracy: 0.9925 - val_loss: 0.0885 - val_accuracy: 0.9724

Epoch 57/250
675/675 [=====] - 7s 11ms/step - loss: 0.0340 - accuracy: 0.9928 - val_loss: 0.0875 - val_accuracy: 0.9725

Epoch 58/250
675/675 [=====] - 7s 10ms/step - loss: 0.0331 - accuracy: 0.9929 - val_loss: 0.0863 - val_accuracy: 0.9729

Epoch 59/250
675/675 [=====] - 7s 11ms/step - loss: 0.0321 - accuracy: 0.9934 - val_loss: 0.0850 - val_accuracy: 0.9729

Epoch 60/250
675/675 [=====] - 6s 9ms/step - loss: 0.0312 - accuracy: 0.9938 - val_loss: 0.0853 - val_accuracy: 0.9738

Epoch 61/250
675/675 [=====] - 8s 11ms/step - loss: 0.0305 - accuracy: 0.9937 - val_loss: 0.0854 - val_accuracy: 0.9731

Epoch 62/250
675/675 [=====] - 6s 9ms/step - loss: 0.0297 - accuracy: 0.9941 - val_loss: 0.0853 - val_accuracy: 0.9729

Epoch 63/250
675/675 [=====] - 8s 12ms/step - loss: 0.0287 - accuracy: 0.9946 - val_loss: 0.0843 - val_accuracy: 0.9731

Epoch 64/250
675/675 [=====] - 6s 9ms/step - loss: 0.0281 - accuracy: 0.9944 - val_loss: 0.0839 - val_accuracy: 0.9728

Epoch 65/250
675/675 [=====] - 8s 12ms/step - loss: 0.0273 - accuracy: 0.9949 - val_loss: 0.0837 - val_accuracy: 0.9735

Epoch 66/250
675/675 [=====] - 6s 9ms/step - loss: 0.0264 - accuracy: 0.9951 - val_loss: 0.0840 - val_accuracy: 0.9737

Epoch 67/250
675/675 [=====] - 8s 11ms/step - loss: 0.0258 - accuracy: 0.9956 - val_loss: 0.0835 - val_accuracy: 0.9739

Epoch 68/250
675/675 [=====] - 6s 9ms/step - loss: 0.0252 - accuracy: 0.9956 - val_loss: 0.0833 - val_accuracy: 0.9730

Epoch 69/250
675/675 [=====] - 8s 11ms/step - loss: 0.0247 - accuracy: 0.9955 - val_loss: 0.0823 - val_accuracy: 0.9736

Epoch 70/250
675/675 [=====] - 6s 9ms/step - loss: 0.0239 - accuracy: 0.9959 - val_loss: 0.0822 - val_accuracy: 0.9744

Epoch 71/250
675/675 [=====] - 7s 11ms/step - loss: 0.0233 - accuracy: 0.9961 - val_loss: 0.0833 - val_accuracy: 0.9732

Epoch 72/250
675/675 [=====] - 6s 10ms/step - loss: 0.0226 -
accuracy: 0.9962 - val_loss: 0.0828 - val_accuracy: 0.9742
Epoch 73/250
675/675 [=====] - 7s 10ms/step - loss: 0.0221 -
accuracy: 0.9966 - val_loss: 0.0820 - val_accuracy: 0.9743
Epoch 74/250
675/675 [=====] - 7s 10ms/step - loss: 0.0214 -
accuracy: 0.9969 - val_loss: 0.0815 - val_accuracy: 0.9744
Epoch 75/250
675/675 [=====] - 7s 10ms/step - loss: 0.0210 -
accuracy: 0.9967 - val_loss: 0.0833 - val_accuracy: 0.9727
Epoch 76/250
675/675 [=====] - 7s 11ms/step - loss: 0.0204 -
accuracy: 0.9970 - val_loss: 0.0824 - val_accuracy: 0.9746
Epoch 77/250
675/675 [=====] - 8s 12ms/step - loss: 0.0199 -
accuracy: 0.9972 - val_loss: 0.0819 - val_accuracy: 0.9742
Epoch 78/250
675/675 [=====] - 7s 11ms/step - loss: 0.0195 -
accuracy: 0.9971 - val_loss: 0.0816 - val_accuracy: 0.9746
Epoch 79/250
675/675 [=====] - 6s 9ms/step - loss: 0.0189 -
accuracy: 0.9976 - val_loss: 0.0824 - val_accuracy: 0.9744
Epoch 80/250
675/675 [=====] - 8s 11ms/step - loss: 0.0185 -
accuracy: 0.9977 - val_loss: 0.0824 - val_accuracy: 0.9745
Epoch 81/250
675/675 [=====] - 6s 9ms/step - loss: 0.0180 -
accuracy: 0.9978 - val_loss: 0.0813 - val_accuracy: 0.9754
Epoch 82/250
675/675 [=====] - 8s 12ms/step - loss: 0.0176 -
accuracy: 0.9978 - val_loss: 0.0816 - val_accuracy: 0.9747
Epoch 83/250
675/675 [=====] - 6s 9ms/step - loss: 0.0172 -
accuracy: 0.9980 - val_loss: 0.0810 - val_accuracy: 0.9744
Epoch 84/250
675/675 [=====] - 8s 12ms/step - loss: 0.0168 -
accuracy: 0.9979 - val_loss: 0.0825 - val_accuracy: 0.9747
Epoch 85/250
675/675 [=====] - 6s 9ms/step - loss: 0.0164 -
accuracy: 0.9982 - val_loss: 0.0804 - val_accuracy: 0.9747
Epoch 86/250
675/675 [=====] - 8s 12ms/step - loss: 0.0160 -
accuracy: 0.9981 - val_loss: 0.0810 - val_accuracy: 0.9749
Epoch 87/250
675/675 [=====] - 6s 9ms/step - loss: 0.0157 -
accuracy: 0.9983 - val_loss: 0.0811 - val_accuracy: 0.9746

Epoch 88/250
675/675 [=====] - 8s 11ms/step - loss: 0.0152 -
accuracy: 0.9982 - val_loss: 0.0807 - val_accuracy: 0.9744

Epoch 89/250
675/675 [=====] - 6s 9ms/step - loss: 0.0149 -
accuracy: 0.9983 - val_loss: 0.0816 - val_accuracy: 0.9744

Epoch 90/250
675/675 [=====] - 7s 11ms/step - loss: 0.0145 -
accuracy: 0.9984 - val_loss: 0.0808 - val_accuracy: 0.9751

Epoch 91/250
675/675 [=====] - 7s 10ms/step - loss: 0.0142 -
accuracy: 0.9985 - val_loss: 0.0805 - val_accuracy: 0.9751

Epoch 92/250
675/675 [=====] - 7s 10ms/step - loss: 0.0139 -
accuracy: 0.9987 - val_loss: 0.0806 - val_accuracy: 0.9746

Epoch 93/250
675/675 [=====] - 7s 10ms/step - loss: 0.0136 -
accuracy: 0.9986 - val_loss: 0.0810 - val_accuracy: 0.9755

Epoch 94/250
675/675 [=====] - 6s 9ms/step - loss: 0.0133 -
accuracy: 0.9987 - val_loss: 0.0807 - val_accuracy: 0.9757

Epoch 95/250
675/675 [=====] - 8s 11ms/step - loss: 0.0129 -
accuracy: 0.9988 - val_loss: 0.0818 - val_accuracy: 0.9759

Epoch 96/250
675/675 [=====] - 6s 9ms/step - loss: 0.0127 -
accuracy: 0.9988 - val_loss: 0.0806 - val_accuracy: 0.9751

Epoch 97/250
675/675 [=====] - 8s 12ms/step - loss: 0.0124 -
accuracy: 0.9987 - val_loss: 0.0807 - val_accuracy: 0.9748

Epoch 98/250
675/675 [=====] - 6s 9ms/step - loss: 0.0121 -
accuracy: 0.9988 - val_loss: 0.0803 - val_accuracy: 0.9756

Epoch 99/250
675/675 [=====] - 8s 12ms/step - loss: 0.0119 -
accuracy: 0.9989 - val_loss: 0.0804 - val_accuracy: 0.9754

Epoch 100/250
675/675 [=====] - 8s 12ms/step - loss: 0.0116 -
accuracy: 0.9989 - val_loss: 0.0811 - val_accuracy: 0.9759

Epoch 101/250
675/675 [=====] - 8s 12ms/step - loss: 0.0114 -
accuracy: 0.9989 - val_loss: 0.0810 - val_accuracy: 0.9753

Epoch 102/250
675/675 [=====] - 6s 9ms/step - loss: 0.0111 -
accuracy: 0.9990 - val_loss: 0.0804 - val_accuracy: 0.9759

Epoch 103/250
675/675 [=====] - 9s 14ms/step - loss: 0.0109 -
accuracy: 0.9991 - val_loss: 0.0810 - val_accuracy: 0.9755

Epoch 104/250
675/675 [=====] - 11s 17ms/step - loss: 0.0106 - accuracy: 0.9991 - val_loss: 0.0808 - val_accuracy: 0.9752

Epoch 105/250
675/675 [=====] - 11s 16ms/step - loss: 0.0104 - accuracy: 0.9991 - val_loss: 0.0806 - val_accuracy: 0.9757

Epoch 106/250
675/675 [=====] - 8s 12ms/step - loss: 0.0102 - accuracy: 0.9991 - val_loss: 0.0808 - val_accuracy: 0.9757

Epoch 107/250
675/675 [=====] - 6s 9ms/step - loss: 0.0100 - accuracy: 0.9993 - val_loss: 0.0809 - val_accuracy: 0.9756

Epoch 108/250
675/675 [=====] - 8s 12ms/step - loss: 0.0098 - accuracy: 0.9993 - val_loss: 0.0812 - val_accuracy: 0.9754

Epoch 109/250
675/675 [=====] - 6s 9ms/step - loss: 0.0096 - accuracy: 0.9992 - val_loss: 0.0807 - val_accuracy: 0.9756

Epoch 110/250
675/675 [=====] - 8s 12ms/step - loss: 0.0093 - accuracy: 0.9993 - val_loss: 0.0808 - val_accuracy: 0.9757

Epoch 111/250
675/675 [=====] - 6s 9ms/step - loss: 0.0092 - accuracy: 0.9994 - val_loss: 0.0805 - val_accuracy: 0.9755

Epoch 112/250
675/675 [=====] - 8s 12ms/step - loss: 0.0091 - accuracy: 0.9993 - val_loss: 0.0808 - val_accuracy: 0.9762

Epoch 113/250
675/675 [=====] - 6s 9ms/step - loss: 0.0089 - accuracy: 0.9993 - val_loss: 0.0808 - val_accuracy: 0.9758

Epoch 114/250
675/675 [=====] - 9s 13ms/step - loss: 0.0087 - accuracy: 0.9994 - val_loss: 0.0808 - val_accuracy: 0.9757

Epoch 115/250
675/675 [=====] - 6s 9ms/step - loss: 0.0085 - accuracy: 0.9995 - val_loss: 0.0814 - val_accuracy: 0.9756

Epoch 116/250
675/675 [=====] - 8s 12ms/step - loss: 0.0083 - accuracy: 0.9995 - val_loss: 0.0813 - val_accuracy: 0.9759

Epoch 117/250
675/675 [=====] - 7s 11ms/step - loss: 0.0082 - accuracy: 0.9995 - val_loss: 0.0812 - val_accuracy: 0.9759

Epoch 118/250
675/675 [=====] - 8s 12ms/step - loss: 0.0080 - accuracy: 0.9995 - val_loss: 0.0808 - val_accuracy: 0.9760

Epoch 119/250
675/675 [=====] - 6s 9ms/step - loss: 0.0079 - accuracy: 0.9995 - val_loss: 0.0810 - val_accuracy: 0.9756

Epoch 120/250
675/675 [=====] - 8s 11ms/step - loss: 0.0077 - accuracy: 0.9996 - val_loss: 0.0808 - val_accuracy: 0.9764

Epoch 121/250
675/675 [=====] - 8s 12ms/step - loss: 0.0076 - accuracy: 0.9995 - val_loss: 0.0812 - val_accuracy: 0.9761

Epoch 122/250
675/675 [=====] - 8s 12ms/step - loss: 0.0075 - accuracy: 0.9995 - val_loss: 0.0811 - val_accuracy: 0.9758

Epoch 123/250
675/675 [=====] - 6s 9ms/step - loss: 0.0073 - accuracy: 0.9996 - val_loss: 0.0813 - val_accuracy: 0.9756

Epoch 124/250
675/675 [=====] - 8s 12ms/step - loss: 0.0072 - accuracy: 0.9996 - val_loss: 0.0810 - val_accuracy: 0.9763

Epoch 125/250
675/675 [=====] - 6s 9ms/step - loss: 0.0070 - accuracy: 0.9997 - val_loss: 0.0814 - val_accuracy: 0.9757

Epoch 126/250
675/675 [=====] - 8s 12ms/step - loss: 0.0069 - accuracy: 0.9997 - val_loss: 0.0817 - val_accuracy: 0.9760

Epoch 127/250
675/675 [=====] - 6s 9ms/step - loss: 0.0068 - accuracy: 0.9997 - val_loss: 0.0814 - val_accuracy: 0.9762

Epoch 128/250
675/675 [=====] - 7s 11ms/step - loss: 0.0067 - accuracy: 0.9997 - val_loss: 0.0815 - val_accuracy: 0.9759

Epoch 129/250
675/675 [=====] - 7s 10ms/step - loss: 0.0066 - accuracy: 0.9997 - val_loss: 0.0818 - val_accuracy: 0.9762

Epoch 130/250
675/675 [=====] - 7s 10ms/step - loss: 0.0064 - accuracy: 0.9997 - val_loss: 0.0817 - val_accuracy: 0.9759

Epoch 131/250
675/675 [=====] - 7s 11ms/step - loss: 0.0063 - accuracy: 0.9997 - val_loss: 0.0819 - val_accuracy: 0.9763

Epoch 132/250
675/675 [=====] - 7s 10ms/step - loss: 0.0062 - accuracy: 0.9998 - val_loss: 0.0820 - val_accuracy: 0.9758

Epoch 133/250
675/675 [=====] - 7s 11ms/step - loss: 0.0061 - accuracy: 0.9998 - val_loss: 0.0821 - val_accuracy: 0.9761

Epoch 134/250
675/675 [=====] - 6s 9ms/step - loss: 0.0060 - accuracy: 0.9997 - val_loss: 0.0821 - val_accuracy: 0.9758

Epoch 135/250
675/675 [=====] - 8s 12ms/step - loss: 0.0059 - accuracy: 0.9997 - val_loss: 0.0817 - val_accuracy: 0.9760

Epoch 136/250
675/675 [=====] - 6s 9ms/step - loss: 0.0058 -
accuracy: 0.9997 - val_loss: 0.0819 - val_accuracy: 0.9761
Epoch 137/250
675/675 [=====] - 8s 12ms/step - loss: 0.0057 -
accuracy: 0.9997 - val_loss: 0.0820 - val_accuracy: 0.9757
Epoch 138/250
675/675 [=====] - 6s 9ms/step - loss: 0.0057 -
accuracy: 0.9998 - val_loss: 0.0818 - val_accuracy: 0.9761
Epoch 139/250
675/675 [=====] - 8s 12ms/step - loss: 0.0056 -
accuracy: 0.9998 - val_loss: 0.0819 - val_accuracy: 0.9765
Epoch 140/250
675/675 [=====] - 6s 9ms/step - loss: 0.0055 -
accuracy: 0.9997 - val_loss: 0.0820 - val_accuracy: 0.9761
Epoch 141/250
675/675 [=====] - 8s 12ms/step - loss: 0.0054 -
accuracy: 0.9998 - val_loss: 0.0820 - val_accuracy: 0.9760
Epoch 142/250
675/675 [=====] - 6s 9ms/step - loss: 0.0053 -
accuracy: 0.9998 - val_loss: 0.0823 - val_accuracy: 0.9761
Epoch 143/250
675/675 [=====] - 8s 12ms/step - loss: 0.0052 -
accuracy: 0.9998 - val_loss: 0.0825 - val_accuracy: 0.9761
Epoch 144/250
675/675 [=====] - 8s 12ms/step - loss: 0.0051 -
accuracy: 0.9999 - val_loss: 0.0826 - val_accuracy: 0.9757
Epoch 145/250
675/675 [=====] - 6s 9ms/step - loss: 0.0050 -
accuracy: 0.9998 - val_loss: 0.0826 - val_accuracy: 0.9760
Epoch 146/250
675/675 [=====] - 8s 12ms/step - loss: 0.0050 -
accuracy: 0.9999 - val_loss: 0.0824 - val_accuracy: 0.9759
Epoch 147/250
675/675 [=====] - 6s 9ms/step - loss: 0.0049 -
accuracy: 0.9999 - val_loss: 0.0826 - val_accuracy: 0.9764
Epoch 148/250
675/675 [=====] - 8s 12ms/step - loss: 0.0048 -
accuracy: 0.9999 - val_loss: 0.0826 - val_accuracy: 0.9761
Epoch 149/250
675/675 [=====] - 6s 9ms/step - loss: 0.0047 -
accuracy: 0.9999 - val_loss: 0.0828 - val_accuracy: 0.9758
Epoch 150/250
675/675 [=====] - 8s 12ms/step - loss: 0.0047 -
accuracy: 0.9999 - val_loss: 0.0829 - val_accuracy: 0.9760
Epoch 151/250
675/675 [=====] - 6s 9ms/step - loss: 0.0046 -
accuracy: 0.9999 - val_loss: 0.0828 - val_accuracy: 0.9761

Epoch 152/250
675/675 [=====] - 8s 12ms/step - loss: 0.0045 - accuracy: 0.9999 - val_loss: 0.0830 - val_accuracy: 0.9756

Epoch 153/250
675/675 [=====] - 6s 9ms/step - loss: 0.0045 - accuracy: 0.9999 - val_loss: 0.0831 - val_accuracy: 0.9760

Epoch 154/250
675/675 [=====] - 8s 12ms/step - loss: 0.0044 - accuracy: 0.9999 - val_loss: 0.0833 - val_accuracy: 0.9761

Epoch 155/250
675/675 [=====] - 6s 9ms/step - loss: 0.0043 - accuracy: 0.9999 - val_loss: 0.0833 - val_accuracy: 0.9758

Epoch 156/250
675/675 [=====] - 8s 12ms/step - loss: 0.0043 - accuracy: 0.9999 - val_loss: 0.0832 - val_accuracy: 0.9759

Epoch 157/250
675/675 [=====] - 6s 9ms/step - loss: 0.0042 - accuracy: 0.9999 - val_loss: 0.0831 - val_accuracy: 0.9763

Epoch 158/250
675/675 [=====] - 8s 11ms/step - loss: 0.0042 - accuracy: 0.9999 - val_loss: 0.0835 - val_accuracy: 0.9760

Epoch 159/250
675/675 [=====] - 6s 10ms/step - loss: 0.0041 - accuracy: 0.9999 - val_loss: 0.0834 - val_accuracy: 0.9761

Epoch 160/250
675/675 [=====] - 7s 10ms/step - loss: 0.0041 - accuracy: 0.9999 - val_loss: 0.0836 - val_accuracy: 0.9762

Epoch 161/250
675/675 [=====] - 7s 10ms/step - loss: 0.0040 - accuracy: 0.9999 - val_loss: 0.0836 - val_accuracy: 0.9761

Epoch 162/250
675/675 [=====] - 7s 10ms/step - loss: 0.0040 - accuracy: 0.9999 - val_loss: 0.0834 - val_accuracy: 0.9759

Epoch 163/250
675/675 [=====] - 7s 11ms/step - loss: 0.0039 - accuracy: 0.9999 - val_loss: 0.0837 - val_accuracy: 0.9761

Epoch 164/250
675/675 [=====] - 7s 10ms/step - loss: 0.0039 - accuracy: 1.0000 - val_loss: 0.0839 - val_accuracy: 0.9761

Epoch 165/250
675/675 [=====] - 9s 14ms/step - loss: 0.0038 - accuracy: 0.9999 - val_loss: 0.0835 - val_accuracy: 0.9756

Epoch 166/250
675/675 [=====] - 7s 10ms/step - loss: 0.0038 - accuracy: 1.0000 - val_loss: 0.0838 - val_accuracy: 0.9761

Epoch 167/250
675/675 [=====] - 7s 11ms/step - loss: 0.0037 - accuracy: 0.9999 - val_loss: 0.0838 - val_accuracy: 0.9763

Epoch 168/250
675/675 [=====] - 7s 10ms/step - loss: 0.0037 - accuracy: 1.0000 - val_loss: 0.0841 - val_accuracy: 0.9760

Epoch 169/250
675/675 [=====] - 7s 11ms/step - loss: 0.0036 - accuracy: 1.0000 - val_loss: 0.0842 - val_accuracy: 0.9763

Epoch 170/250
675/675 [=====] - 6s 9ms/step - loss: 0.0036 - accuracy: 1.0000 - val_loss: 0.0845 - val_accuracy: 0.9756

Epoch 171/250
675/675 [=====] - 8s 12ms/step - loss: 0.0035 - accuracy: 1.0000 - val_loss: 0.0844 - val_accuracy: 0.9763

Epoch 172/250
675/675 [=====] - 6s 9ms/step - loss: 0.0035 - accuracy: 1.0000 - val_loss: 0.0841 - val_accuracy: 0.9760

Epoch 173/250
675/675 [=====] - 8s 12ms/step - loss: 0.0034 - accuracy: 1.0000 - val_loss: 0.0845 - val_accuracy: 0.9757

Epoch 174/250
675/675 [=====] - 6s 9ms/step - loss: 0.0034 - accuracy: 1.0000 - val_loss: 0.0846 - val_accuracy: 0.9762

Epoch 175/250
675/675 [=====] - 8s 12ms/step - loss: 0.0033 - accuracy: 1.0000 - val_loss: 0.0848 - val_accuracy: 0.9761

Epoch 176/250
675/675 [=====] - 6s 9ms/step - loss: 0.0033 - accuracy: 1.0000 - val_loss: 0.0844 - val_accuracy: 0.9757

Epoch 177/250
675/675 [=====] - 8s 12ms/step - loss: 0.0033 - accuracy: 1.0000 - val_loss: 0.0847 - val_accuracy: 0.9761

Epoch 178/250
675/675 [=====] - 6s 9ms/step - loss: 0.0032 - accuracy: 1.0000 - val_loss: 0.0846 - val_accuracy: 0.9761

Epoch 179/250
675/675 [=====] - 8s 12ms/step - loss: 0.0032 - accuracy: 1.0000 - val_loss: 0.0845 - val_accuracy: 0.9761

Epoch 180/250
675/675 [=====] - 6s 9ms/step - loss: 0.0032 - accuracy: 1.0000 - val_loss: 0.0849 - val_accuracy: 0.9760

Epoch 181/250
675/675 [=====] - 8s 13ms/step - loss: 0.0031 - accuracy: 1.0000 - val_loss: 0.0848 - val_accuracy: 0.9763

Epoch 182/250
675/675 [=====] - 6s 9ms/step - loss: 0.0031 - accuracy: 1.0000 - val_loss: 0.0849 - val_accuracy: 0.9757

Epoch 183/250
675/675 [=====] - 8s 12ms/step - loss: 0.0031 - accuracy: 1.0000 - val_loss: 0.0850 - val_accuracy: 0.9757

Epoch 184/250
675/675 [=====] - 6s 9ms/step - loss: 0.0030 - accuracy: 1.0000 - val_loss: 0.0851 - val_accuracy: 0.9758

Epoch 185/250
675/675 [=====] - 7s 11ms/step - loss: 0.0030 - accuracy: 1.0000 - val_loss: 0.0851 - val_accuracy: 0.9760

Epoch 186/250
675/675 [=====] - 8s 13ms/step - loss: 0.0030 - accuracy: 1.0000 - val_loss: 0.0850 - val_accuracy: 0.9761

Epoch 187/250
675/675 [=====] - 8s 11ms/step - loss: 0.0029 - accuracy: 1.0000 - val_loss: 0.0851 - val_accuracy: 0.9761

Epoch 188/250
675/675 [=====] - 7s 10ms/step - loss: 0.0029 - accuracy: 1.0000 - val_loss: 0.0850 - val_accuracy: 0.9761

Epoch 189/250
675/675 [=====] - 8s 12ms/step - loss: 0.0029 - accuracy: 1.0000 - val_loss: 0.0855 - val_accuracy: 0.9761

Epoch 190/250
675/675 [=====] - 7s 10ms/step - loss: 0.0028 - accuracy: 1.0000 - val_loss: 0.0853 - val_accuracy: 0.9761

Epoch 191/250
675/675 [=====] - 7s 11ms/step - loss: 0.0028 - accuracy: 1.0000 - val_loss: 0.0853 - val_accuracy: 0.9758

Epoch 192/250
675/675 [=====] - 7s 10ms/step - loss: 0.0028 - accuracy: 1.0000 - val_loss: 0.0854 - val_accuracy: 0.9763

Epoch 193/250
675/675 [=====] - 7s 10ms/step - loss: 0.0028 - accuracy: 1.0000 - val_loss: 0.0857 - val_accuracy: 0.9757

Epoch 194/250
675/675 [=====] - 7s 11ms/step - loss: 0.0027 - accuracy: 1.0000 - val_loss: 0.0856 - val_accuracy: 0.9761

Epoch 195/250
675/675 [=====] - 7s 10ms/step - loss: 0.0027 - accuracy: 1.0000 - val_loss: 0.0856 - val_accuracy: 0.9761

Epoch 196/250
675/675 [=====] - 7s 11ms/step - loss: 0.0027 - accuracy: 1.0000 - val_loss: 0.0857 - val_accuracy: 0.9762

Epoch 197/250
675/675 [=====] - 7s 10ms/step - loss: 0.0026 - accuracy: 1.0000 - val_loss: 0.0856 - val_accuracy: 0.9765

Epoch 198/250
675/675 [=====] - 8s 11ms/step - loss: 0.0026 - accuracy: 1.0000 - val_loss: 0.0859 - val_accuracy: 0.9764

Epoch 199/250
675/675 [=====] - 7s 10ms/step - loss: 0.0026 - accuracy: 1.0000 - val_loss: 0.0859 - val_accuracy: 0.9763

Epoch 200/250
675/675 [=====] - 7s 11ms/step - loss: 0.0026 - accuracy: 1.0000 - val_loss: 0.0860 - val_accuracy: 0.9758

Epoch 201/250
675/675 [=====] - 6s 9ms/step - loss: 0.0025 - accuracy: 1.0000 - val_loss: 0.0861 - val_accuracy: 0.9759

Epoch 202/250
675/675 [=====] - 8s 12ms/step - loss: 0.0025 - accuracy: 1.0000 - val_loss: 0.0859 - val_accuracy: 0.9768

Epoch 203/250
675/675 [=====] - 6s 9ms/step - loss: 0.0025 - accuracy: 1.0000 - val_loss: 0.0863 - val_accuracy: 0.9761

Epoch 204/250
675/675 [=====] - 8s 12ms/step - loss: 0.0025 - accuracy: 1.0000 - val_loss: 0.0861 - val_accuracy: 0.9760

Epoch 205/250
675/675 [=====] - 6s 9ms/step - loss: 0.0024 - accuracy: 1.0000 - val_loss: 0.0862 - val_accuracy: 0.9765

Epoch 206/250
675/675 [=====] - 8s 12ms/step - loss: 0.0024 - accuracy: 1.0000 - val_loss: 0.0865 - val_accuracy: 0.9765

Epoch 207/250
675/675 [=====] - 8s 12ms/step - loss: 0.0024 - accuracy: 1.0000 - val_loss: 0.0863 - val_accuracy: 0.9762

Epoch 208/250
675/675 [=====] - 8s 12ms/step - loss: 0.0024 - accuracy: 1.0000 - val_loss: 0.0863 - val_accuracy: 0.9762

Epoch 209/250
675/675 [=====] - 6s 9ms/step - loss: 0.0024 - accuracy: 1.0000 - val_loss: 0.0863 - val_accuracy: 0.9761

Epoch 210/250
675/675 [=====] - 8s 12ms/step - loss: 0.0023 - accuracy: 1.0000 - val_loss: 0.0864 - val_accuracy: 0.9761

Epoch 211/250
675/675 [=====] - 6s 9ms/step - loss: 0.0023 - accuracy: 1.0000 - val_loss: 0.0865 - val_accuracy: 0.9765

Epoch 212/250
675/675 [=====] - 8s 12ms/step - loss: 0.0023 - accuracy: 1.0000 - val_loss: 0.0866 - val_accuracy: 0.9763

Epoch 213/250
675/675 [=====] - 6s 9ms/step - loss: 0.0023 - accuracy: 1.0000 - val_loss: 0.0868 - val_accuracy: 0.9763

Epoch 214/250
675/675 [=====] - 8s 12ms/step - loss: 0.0023 - accuracy: 1.0000 - val_loss: 0.0866 - val_accuracy: 0.9766

Epoch 215/250
675/675 [=====] - 6s 9ms/step - loss: 0.0022 - accuracy: 1.0000 - val_loss: 0.0867 - val_accuracy: 0.9763

Epoch 216/250
675/675 [=====] - 8s 11ms/step - loss: 0.0022 - accuracy: 1.0000 - val_loss: 0.0867 - val_accuracy: 0.9764
Epoch 217/250
675/675 [=====] - 7s 10ms/step - loss: 0.0022 - accuracy: 1.0000 - val_loss: 0.0868 - val_accuracy: 0.9763
Epoch 218/250
675/675 [=====] - 7s 11ms/step - loss: 0.0022 - accuracy: 1.0000 - val_loss: 0.0871 - val_accuracy: 0.9762
Epoch 219/250
675/675 [=====] - 7s 10ms/step - loss: 0.0022 - accuracy: 1.0000 - val_loss: 0.0872 - val_accuracy: 0.9764
Epoch 220/250
675/675 [=====] - 7s 10ms/step - loss: 0.0021 - accuracy: 1.0000 - val_loss: 0.0872 - val_accuracy: 0.9763
Epoch 221/250
675/675 [=====] - 7s 11ms/step - loss: 0.0021 - accuracy: 1.0000 - val_loss: 0.0870 - val_accuracy: 0.9764
Epoch 222/250
675/675 [=====] - 7s 10ms/step - loss: 0.0021 - accuracy: 1.0000 - val_loss: 0.0870 - val_accuracy: 0.9764
Epoch 223/250
675/675 [=====] - 8s 11ms/step - loss: 0.0021 - accuracy: 1.0000 - val_loss: 0.0872 - val_accuracy: 0.9767
Epoch 224/250
675/675 [=====] - 6s 9ms/step - loss: 0.0021 - accuracy: 1.0000 - val_loss: 0.0872 - val_accuracy: 0.9766
Epoch 225/250
675/675 [=====] - 8s 12ms/step - loss: 0.0020 - accuracy: 1.0000 - val_loss: 0.0874 - val_accuracy: 0.9769
Epoch 226/250
675/675 [=====] - 6s 9ms/step - loss: 0.0020 - accuracy: 1.0000 - val_loss: 0.0876 - val_accuracy: 0.9760
Epoch 227/250
675/675 [=====] - 10s 15ms/step - loss: 0.0020 - accuracy: 1.0000 - val_loss: 0.0874 - val_accuracy: 0.9765
Epoch 228/250
675/675 [=====] - 6s 10ms/step - loss: 0.0020 - accuracy: 1.0000 - val_loss: 0.0875 - val_accuracy: 0.9763
Epoch 229/250
675/675 [=====] - 8s 12ms/step - loss: 0.0020 - accuracy: 1.0000 - val_loss: 0.0875 - val_accuracy: 0.9765
Epoch 230/250
675/675 [=====] - 6s 9ms/step - loss: 0.0020 - accuracy: 1.0000 - val_loss: 0.0875 - val_accuracy: 0.9764
Epoch 231/250
675/675 [=====] - 8s 12ms/step - loss: 0.0020 - accuracy: 1.0000 - val_loss: 0.0876 - val_accuracy: 0.9764

Epoch 232/250
675/675 [=====] - 6s 9ms/step - loss: 0.0019 - accuracy: 1.0000 - val_loss: 0.0877 - val_accuracy: 0.9765

Epoch 233/250
675/675 [=====] - 8s 11ms/step - loss: 0.0019 - accuracy: 1.0000 - val_loss: 0.0879 - val_accuracy: 0.9764

Epoch 234/250
675/675 [=====] - 7s 10ms/step - loss: 0.0019 - accuracy: 1.0000 - val_loss: 0.0877 - val_accuracy: 0.9764

Epoch 235/250
675/675 [=====] - 8s 11ms/step - loss: 0.0019 - accuracy: 1.0000 - val_loss: 0.0878 - val_accuracy: 0.9763

Epoch 236/250
675/675 [=====] - 7s 10ms/step - loss: 0.0019 - accuracy: 1.0000 - val_loss: 0.0879 - val_accuracy: 0.9764

Epoch 237/250
675/675 [=====] - 7s 11ms/step - loss: 0.0019 - accuracy: 1.0000 - val_loss: 0.0879 - val_accuracy: 0.9763

Epoch 238/250
675/675 [=====] - 7s 10ms/step - loss: 0.0018 - accuracy: 1.0000 - val_loss: 0.0880 - val_accuracy: 0.9765

Epoch 239/250
675/675 [=====] - 8s 12ms/step - loss: 0.0018 - accuracy: 1.0000 - val_loss: 0.0881 - val_accuracy: 0.9765

Epoch 240/250
675/675 [=====] - 7s 10ms/step - loss: 0.0018 - accuracy: 1.0000 - val_loss: 0.0879 - val_accuracy: 0.9767

Epoch 241/250
675/675 [=====] - 7s 10ms/step - loss: 0.0018 - accuracy: 1.0000 - val_loss: 0.0881 - val_accuracy: 0.9765

Epoch 242/250
675/675 [=====] - 8s 11ms/step - loss: 0.0018 - accuracy: 1.0000 - val_loss: 0.0882 - val_accuracy: 0.9769

Epoch 243/250
675/675 [=====] - 6s 9ms/step - loss: 0.0018 - accuracy: 1.0000 - val_loss: 0.0883 - val_accuracy: 0.9765

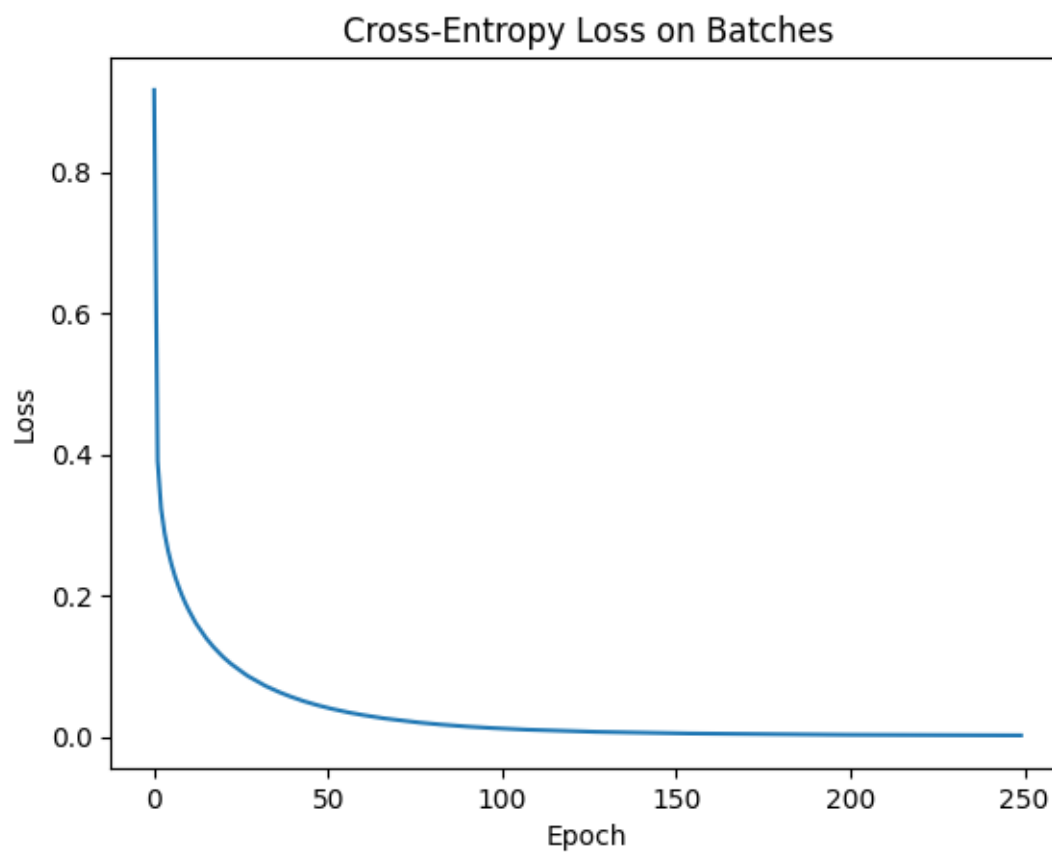
Epoch 244/250
675/675 [=====] - 8s 12ms/step - loss: 0.0018 - accuracy: 1.0000 - val_loss: 0.0883 - val_accuracy: 0.9763

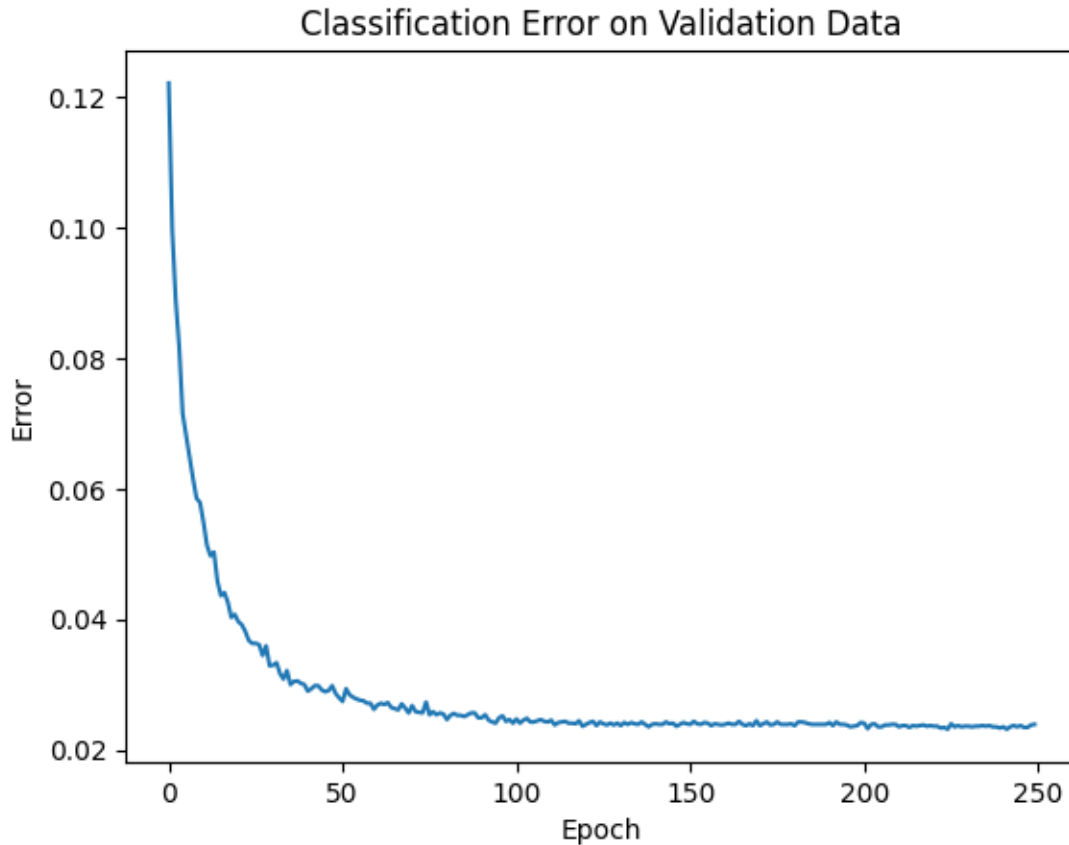
Epoch 245/250
675/675 [=====] - 6s 9ms/step - loss: 0.0018 - accuracy: 1.0000 - val_loss: 0.0883 - val_accuracy: 0.9765

Epoch 246/250
675/675 [=====] - 8s 12ms/step - loss: 0.0017 - accuracy: 1.0000 - val_loss: 0.0884 - val_accuracy: 0.9763

Epoch 247/250
675/675 [=====] - 8s 12ms/step - loss: 0.0017 - accuracy: 1.0000 - val_loss: 0.0885 - val_accuracy: 0.9766

Epoch 248/250
675/675 [=====] - 8s 12ms/step - loss: 0.0017 -
accuracy: 1.0000 - val_loss: 0.0885 - val_accuracy: 0.9766
Epoch 249/250
675/675 [=====] - 6s 9ms/step - loss: 0.0017 -
accuracy: 1.0000 - val_loss: 0.0888 - val_accuracy: 0.9762
Epoch 250/250
675/675 [=====] - 8s 12ms/step - loss: 0.0017 -
accuracy: 1.0000 - val_loss: 0.0885 - val_accuracy: 0.9761





Test Accuracy: 0.9799000024795532

In this experiment, we constructed a three-layer neural network with a higher number of neurons in each layer compared to previous models. We also increased the training duration to 250 epochs.

After training the model, we obtained 100% accuracy while training and test accuracy of 97.99%, indicating that it was able to perfectly fit the training data. This exceptional performance suggests that the model had sufficient capacity to capture the complex patterns present in the training dataset. Also this 100% accuracy was achievable due to simplicity of the data and actually it was achieved by 170 epochs only.

Furthermore, the model achieved an accuracy of 97.6% on the validation set. This result demonstrates that the model generalized well to unseen data, as it performed with high accuracy on samples it had not been exposed to during training. This validates the model's ability to capture the underlying patterns in the data and make accurate predictions.

For the purpose of comparison, we employed the SGD optimizer with a learning rate of 0.01 and a batch size of 64, consistent with the previous questions. This allows us to evaluate the impact of the increased model complexity and training duration while keeping other factors constant.

4 Que 4:

Repeat Question #3 but train the network with the following regularizations:

L2-norm, dropout, and early-stopping. Compare with the results of Question #3.

Comments: No need to implement them by your own (this is true in general); the software framework (ie, the DL platform) typically provides implementations for all the regularizations methods discussed in class. Early stopping is done in training, so you only need to tune your training code slightly.

```
[ ]: from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping

# Define the three-layer feedforward neural network with regularizations
model = Sequential()
model.add(Dense(500, activation='relu', input_shape=(784,),
    ↪kernel_regularizer='l2'))
model.add(Dropout(0.5))
model.add(Dense(500, activation='relu', kernel_regularizer='l2'))
model.add(Dropout(0.5))
model.add(Dense(10, activation='sigmoid'))

# Compile the model
model.compile(loss='categorical_crossentropy', optimizer=SGD(learning_rate=0.
    ↪0.01), metrics=['accuracy'])

# Define early stopping callback
early_stopping = EarlyStopping(monitor='val_loss', patience=10)

# Train the model with regularizations and early stopping
history = model.fit(x_train, y_train, batch_size=64, epochs=250,
    ↪validation_split=0.2, callbacks=[early_stopping], verbose=1)

# Evaluate the model on the testing set
_, test_accuracy = model.evaluate(x_test, y_test, verbose=0)

print("\n")
# Plot the cross-entropy loss on the batches
plt.plot(history.history['loss'])
plt.title('Cross-Entropy Loss on Batches (with Regularizations)')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.show()
print("\n")

# Plot the classification error on the validation data
plt.plot(1 - np.array(history.history['val_accuracy']))
```

```
plt.title('Classification Error on Validation Data (with Regularizations)')
plt.xlabel('Epoch')
plt.ylabel('Error')
plt.show()

print("\n\n")
# Print the test accuracy
print(f'Test Accuracy: {test_accuracy}')
```

```
Epoch 1/250
675/675 [=====] - 8s 11ms/step - loss: 11.0878 -
accuracy: 0.5968 - val_loss: 9.1323 - val_accuracy: 0.8585
Epoch 2/250
675/675 [=====] - 9s 13ms/step - loss: 8.1900 -
accuracy: 0.8062 - val_loss: 7.0144 - val_accuracy: 0.8893
Epoch 3/250
675/675 [=====] - 7s 11ms/step - loss: 6.3363 -
accuracy: 0.8498 - val_loss: 5.4540 - val_accuracy: 0.9004
Epoch 4/250
675/675 [=====] - 9s 13ms/step - loss: 4.9533 -
accuracy: 0.8682 - val_loss: 4.2735 - val_accuracy: 0.9064
Epoch 5/250
675/675 [=====] - 8s 12ms/step - loss: 3.9040 -
accuracy: 0.8798 - val_loss: 3.3728 - val_accuracy: 0.9111
Epoch 6/250
675/675 [=====] - 8s 12ms/step - loss: 3.1057 -
accuracy: 0.8882 - val_loss: 2.6870 - val_accuracy: 0.9156
Epoch 7/250
675/675 [=====] - 9s 14ms/step - loss: 2.4973 -
accuracy: 0.8929 - val_loss: 2.1638 - val_accuracy: 0.9194
Epoch 8/250
675/675 [=====] - 7s 11ms/step - loss: 2.0300 -
accuracy: 0.8994 - val_loss: 1.7635 - val_accuracy: 0.9217
Epoch 9/250
675/675 [=====] - 9s 13ms/step - loss: 1.6749 -
accuracy: 0.9025 - val_loss: 1.4568 - val_accuracy: 0.9227
Epoch 10/250
675/675 [=====] - 8s 12ms/step - loss: 1.3998 -
accuracy: 0.9066 - val_loss: 1.2230 - val_accuracy: 0.9253
Epoch 11/250
675/675 [=====] - 10s 15ms/step - loss: 1.1903 -
accuracy: 0.9090 - val_loss: 1.0430 - val_accuracy: 0.9273
Epoch 12/250
675/675 [=====] - 7s 11ms/step - loss: 1.0329 -
accuracy: 0.9110 - val_loss: 0.9052 - val_accuracy: 0.9297
Epoch 13/250
675/675 [=====] - 9s 13ms/step - loss: 0.9067 -
```

accuracy: 0.9140 - val_loss: 0.7981 - val_accuracy: 0.9324
Epoch 14/250
675/675 [=====] - 9s 13ms/step - loss: 0.8137 -
accuracy: 0.9141 - val_loss: 0.7166 - val_accuracy: 0.9319
Epoch 15/250
675/675 [=====] - 8s 12ms/step - loss: 0.7368 -
accuracy: 0.9184 - val_loss: 0.6525 - val_accuracy: 0.9340
Epoch 16/250
675/675 [=====] - 9s 13ms/step - loss: 0.6809 -
accuracy: 0.9182 - val_loss: 0.6043 - val_accuracy: 0.9340
Epoch 17/250
675/675 [=====] - 7s 11ms/step - loss: 0.6333 -
accuracy: 0.9211 - val_loss: 0.5639 - val_accuracy: 0.9369
Epoch 18/250
675/675 [=====] - 9s 14ms/step - loss: 0.5972 -
accuracy: 0.9223 - val_loss: 0.5338 - val_accuracy: 0.9387
Epoch 19/250
675/675 [=====] - 7s 11ms/step - loss: 0.5680 -
accuracy: 0.9248 - val_loss: 0.5128 - val_accuracy: 0.9344
Epoch 20/250
675/675 [=====] - 9s 14ms/step - loss: 0.5458 -
accuracy: 0.9254 - val_loss: 0.4891 - val_accuracy: 0.9395
Epoch 21/250
675/675 [=====] - 9s 13ms/step - loss: 0.5283 -
accuracy: 0.9258 - val_loss: 0.4750 - val_accuracy: 0.9396
Epoch 22/250
675/675 [=====] - 8s 12ms/step - loss: 0.5142 -
accuracy: 0.9266 - val_loss: 0.4614 - val_accuracy: 0.9409
Epoch 23/250
675/675 [=====] - 9s 13ms/step - loss: 0.5028 -
accuracy: 0.9277 - val_loss: 0.4509 - val_accuracy: 0.9408
Epoch 24/250
675/675 [=====] - 7s 11ms/step - loss: 0.4897 -
accuracy: 0.9293 - val_loss: 0.4420 - val_accuracy: 0.9418
Epoch 25/250
675/675 [=====] - 9s 13ms/step - loss: 0.4815 -
accuracy: 0.9306 - val_loss: 0.4334 - val_accuracy: 0.9444
Epoch 26/250
675/675 [=====] - 7s 11ms/step - loss: 0.4734 -
accuracy: 0.9314 - val_loss: 0.4277 - val_accuracy: 0.9439
Epoch 27/250
675/675 [=====] - 10s 15ms/step - loss: 0.4684 -
accuracy: 0.9311 - val_loss: 0.4221 - val_accuracy: 0.9449
Epoch 28/250
675/675 [=====] - 10s 14ms/step - loss: 0.4641 -
accuracy: 0.9312 - val_loss: 0.4172 - val_accuracy: 0.9445
Epoch 29/250
675/675 [=====] - 8s 11ms/step - loss: 0.4584 -

accuracy: 0.9322 - val_loss: 0.4123 - val_accuracy: 0.9452
Epoch 30/250
675/675 [=====] - 9s 13ms/step - loss: 0.4519 -
accuracy: 0.9337 - val_loss: 0.4084 - val_accuracy: 0.9464
Epoch 31/250
675/675 [=====] - 7s 11ms/step - loss: 0.4487 -
accuracy: 0.9337 - val_loss: 0.4058 - val_accuracy: 0.9463
Epoch 32/250
675/675 [=====] - 9s 13ms/step - loss: 0.4439 -
accuracy: 0.9346 - val_loss: 0.4018 - val_accuracy: 0.9463
Epoch 33/250
675/675 [=====] - 8s 12ms/step - loss: 0.4433 -
accuracy: 0.9354 - val_loss: 0.3992 - val_accuracy: 0.9476
Epoch 34/250
675/675 [=====] - 8s 12ms/step - loss: 0.4372 -
accuracy: 0.9362 - val_loss: 0.3953 - val_accuracy: 0.9469
Epoch 35/250
675/675 [=====] - 9s 13ms/step - loss: 0.4331 -
accuracy: 0.9377 - val_loss: 0.3931 - val_accuracy: 0.9480
Epoch 36/250
675/675 [=====] - 8s 11ms/step - loss: 0.4321 -
accuracy: 0.9372 - val_loss: 0.3931 - val_accuracy: 0.9469
Epoch 37/250
675/675 [=====] - 9s 13ms/step - loss: 0.4276 -
accuracy: 0.9382 - val_loss: 0.3895 - val_accuracy: 0.9473
Epoch 38/250
675/675 [=====] - 7s 11ms/step - loss: 0.4264 -
accuracy: 0.9381 - val_loss: 0.3849 - val_accuracy: 0.9508
Epoch 39/250
675/675 [=====] - 9s 14ms/step - loss: 0.4242 -
accuracy: 0.9383 - val_loss: 0.3838 - val_accuracy: 0.9501
Epoch 40/250
675/675 [=====] - 7s 11ms/step - loss: 0.4197 -
accuracy: 0.9392 - val_loss: 0.3814 - val_accuracy: 0.9511
Epoch 41/250
675/675 [=====] - 9s 13ms/step - loss: 0.4168 -
accuracy: 0.9398 - val_loss: 0.3792 - val_accuracy: 0.9503
Epoch 42/250
675/675 [=====] - 8s 12ms/step - loss: 0.4174 -
accuracy: 0.9406 - val_loss: 0.3778 - val_accuracy: 0.9503
Epoch 43/250
675/675 [=====] - 8s 12ms/step - loss: 0.4152 -
accuracy: 0.9401 - val_loss: 0.3771 - val_accuracy: 0.9497
Epoch 44/250
675/675 [=====] - 10s 15ms/step - loss: 0.4125 -
accuracy: 0.9413 - val_loss: 0.3733 - val_accuracy: 0.9517
Epoch 45/250
675/675 [=====] - 9s 13ms/step - loss: 0.4099 -

accuracy: 0.9417 - val_loss: 0.3720 - val_accuracy: 0.9503
Epoch 46/250
675/675 [=====] - 9s 13ms/step - loss: 0.4087 -
accuracy: 0.9414 - val_loss: 0.3699 - val_accuracy: 0.9525
Epoch 47/250
675/675 [=====] - 7s 11ms/step - loss: 0.4062 -
accuracy: 0.9423 - val_loss: 0.3673 - val_accuracy: 0.9515
Epoch 48/250
675/675 [=====] - 9s 13ms/step - loss: 0.4031 -
accuracy: 0.9436 - val_loss: 0.3670 - val_accuracy: 0.9512
Epoch 49/250
675/675 [=====] - 8s 12ms/step - loss: 0.4026 -
accuracy: 0.9431 - val_loss: 0.3651 - val_accuracy: 0.9524
Epoch 50/250
675/675 [=====] - 8s 12ms/step - loss: 0.3986 -
accuracy: 0.9443 - val_loss: 0.3644 - val_accuracy: 0.9519
Epoch 51/250
675/675 [=====] - 9s 14ms/step - loss: 0.3991 -
accuracy: 0.9433 - val_loss: 0.3628 - val_accuracy: 0.9527
Epoch 52/250
675/675 [=====] - 7s 11ms/step - loss: 0.3972 -
accuracy: 0.9444 - val_loss: 0.3617 - val_accuracy: 0.9519
Epoch 53/250
675/675 [=====] - 12s 17ms/step - loss: 0.3963 -
accuracy: 0.9440 - val_loss: 0.3591 - val_accuracy: 0.9536
Epoch 54/250
675/675 [=====] - 9s 14ms/step - loss: 0.3937 -
accuracy: 0.9445 - val_loss: 0.3579 - val_accuracy: 0.9542
Epoch 55/250
675/675 [=====] - 7s 11ms/step - loss: 0.3915 -
accuracy: 0.9452 - val_loss: 0.3556 - val_accuracy: 0.9544
Epoch 56/250
675/675 [=====] - 9s 14ms/step - loss: 0.3895 -
accuracy: 0.9454 - val_loss: 0.3558 - val_accuracy: 0.9526
Epoch 57/250
675/675 [=====] - 8s 11ms/step - loss: 0.3901 -
accuracy: 0.9443 - val_loss: 0.3538 - val_accuracy: 0.9538
Epoch 58/250
675/675 [=====] - 9s 14ms/step - loss: 0.3877 -
accuracy: 0.9460 - val_loss: 0.3522 - val_accuracy: 0.9551
Epoch 59/250
675/675 [=====] - 8s 12ms/step - loss: 0.3865 -
accuracy: 0.9455 - val_loss: 0.3507 - val_accuracy: 0.9551
Epoch 60/250
675/675 [=====] - 9s 13ms/step - loss: 0.3871 -
accuracy: 0.9455 - val_loss: 0.3509 - val_accuracy: 0.9533
Epoch 61/250
675/675 [=====] - 10s 15ms/step - loss: 0.3845 -

accuracy: 0.9471 - val_loss: 0.3498 - val_accuracy: 0.9553
 Epoch 62/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3825 -
 accuracy: 0.9472 - val_loss: 0.3484 - val_accuracy: 0.9556
 Epoch 63/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3810 -
 accuracy: 0.9468 - val_loss: 0.3466 - val_accuracy: 0.9543
 Epoch 64/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3791 -
 accuracy: 0.9473 - val_loss: 0.3470 - val_accuracy: 0.9542
 Epoch 65/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3799 -
 accuracy: 0.9457 - val_loss: 0.3454 - val_accuracy: 0.9556
 Epoch 66/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3765 -
 accuracy: 0.9474 - val_loss: 0.3439 - val_accuracy: 0.9545
 Epoch 67/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3769 -
 accuracy: 0.9462 - val_loss: 0.3421 - val_accuracy: 0.9556
 Epoch 68/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3747 -
 accuracy: 0.9480 - val_loss: 0.3414 - val_accuracy: 0.9557
 Epoch 69/250
 675/675 [=====] - 7s 11ms/step - loss: 0.3747 -
 accuracy: 0.9474 - val_loss: 0.3404 - val_accuracy: 0.9567
 Epoch 70/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3748 -
 accuracy: 0.9483 - val_loss: 0.3399 - val_accuracy: 0.9557
 Epoch 71/250
 675/675 [=====] - 8s 11ms/step - loss: 0.3710 -
 accuracy: 0.9479 - val_loss: 0.3387 - val_accuracy: 0.9573
 Epoch 72/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3721 -
 accuracy: 0.9481 - val_loss: 0.3388 - val_accuracy: 0.9567
 Epoch 73/250
 675/675 [=====] - 8s 11ms/step - loss: 0.3699 -
 accuracy: 0.9488 - val_loss: 0.3374 - val_accuracy: 0.9555
 Epoch 74/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3696 -
 accuracy: 0.9489 - val_loss: 0.3358 - val_accuracy: 0.9567
 Epoch 75/250
 675/675 [=====] - 8s 13ms/step - loss: 0.3680 -
 accuracy: 0.9485 - val_loss: 0.3357 - val_accuracy: 0.9570
 Epoch 76/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3682 -
 accuracy: 0.9484 - val_loss: 0.3349 - val_accuracy: 0.9561
 Epoch 77/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3668 -

accuracy: 0.9498 - val_loss: 0.3343 - val_accuracy: 0.9562
Epoch 78/250
675/675 [=====] - 9s 14ms/step - loss: 0.3660 -
accuracy: 0.9489 - val_loss: 0.3313 - val_accuracy: 0.9570
Epoch 79/250
675/675 [=====] - 9s 14ms/step - loss: 0.3637 -
accuracy: 0.9497 - val_loss: 0.3312 - val_accuracy: 0.9575
Epoch 80/250
675/675 [=====] - 9s 13ms/step - loss: 0.3622 -
accuracy: 0.9498 - val_loss: 0.3307 - val_accuracy: 0.9580
Epoch 81/250
675/675 [=====] - 8s 11ms/step - loss: 0.3616 -
accuracy: 0.9503 - val_loss: 0.3304 - val_accuracy: 0.9579
Epoch 82/250
675/675 [=====] - 9s 13ms/step - loss: 0.3619 -
accuracy: 0.9503 - val_loss: 0.3299 - val_accuracy: 0.9570
Epoch 83/250
675/675 [=====] - 7s 11ms/step - loss: 0.3627 -
accuracy: 0.9494 - val_loss: 0.3288 - val_accuracy: 0.9582
Epoch 84/250
675/675 [=====] - 9s 13ms/step - loss: 0.3605 -
accuracy: 0.9498 - val_loss: 0.3283 - val_accuracy: 0.9576
Epoch 85/250
675/675 [=====] - 8s 12ms/step - loss: 0.3591 -
accuracy: 0.9499 - val_loss: 0.3294 - val_accuracy: 0.9566
Epoch 86/250
675/675 [=====] - 9s 13ms/step - loss: 0.3578 -
accuracy: 0.9515 - val_loss: 0.3263 - val_accuracy: 0.9583
Epoch 87/250
675/675 [=====] - 9s 14ms/step - loss: 0.3564 -
accuracy: 0.9498 - val_loss: 0.3253 - val_accuracy: 0.9587
Epoch 88/250
675/675 [=====] - 8s 11ms/step - loss: 0.3581 -
accuracy: 0.9499 - val_loss: 0.3244 - val_accuracy: 0.9576
Epoch 89/250
675/675 [=====] - 9s 13ms/step - loss: 0.3562 -
accuracy: 0.9508 - val_loss: 0.3237 - val_accuracy: 0.9574
Epoch 90/250
675/675 [=====] - 7s 11ms/step - loss: 0.3536 -
accuracy: 0.9509 - val_loss: 0.3245 - val_accuracy: 0.9589
Epoch 91/250
675/675 [=====] - 9s 14ms/step - loss: 0.3537 -
accuracy: 0.9518 - val_loss: 0.3229 - val_accuracy: 0.9581
Epoch 92/250
675/675 [=====] - 8s 12ms/step - loss: 0.3537 -
accuracy: 0.9516 - val_loss: 0.3232 - val_accuracy: 0.9585
Epoch 93/250
675/675 [=====] - 8s 13ms/step - loss: 0.3530 -

accuracy: 0.9518 - val_loss: 0.3212 - val_accuracy: 0.9590
 Epoch 94/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3518 -
 accuracy: 0.9511 - val_loss: 0.3204 - val_accuracy: 0.9587
 Epoch 95/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3520 -
 accuracy: 0.9509 - val_loss: 0.3200 - val_accuracy: 0.9593
 Epoch 96/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3521 -
 accuracy: 0.9513 - val_loss: 0.3183 - val_accuracy: 0.9585
 Epoch 97/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3472 -
 accuracy: 0.9532 - val_loss: 0.3197 - val_accuracy: 0.9581
 Epoch 98/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3516 -
 accuracy: 0.9515 - val_loss: 0.3179 - val_accuracy: 0.9591
 Epoch 99/250
 675/675 [=====] - 10s 14ms/step - loss: 0.3473 -
 accuracy: 0.9518 - val_loss: 0.3173 - val_accuracy: 0.9593
 Epoch 100/250
 675/675 [=====] - 7s 11ms/step - loss: 0.3499 -
 accuracy: 0.9511 - val_loss: 0.3171 - val_accuracy: 0.9602
 Epoch 101/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3462 -
 accuracy: 0.9525 - val_loss: 0.3160 - val_accuracy: 0.9589
 Epoch 102/250
 675/675 [=====] - 7s 11ms/step - loss: 0.3469 -
 accuracy: 0.9518 - val_loss: 0.3159 - val_accuracy: 0.9593
 Epoch 103/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3458 -
 accuracy: 0.9525 - val_loss: 0.3138 - val_accuracy: 0.9600
 Epoch 104/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3453 -
 accuracy: 0.9517 - val_loss: 0.3205 - val_accuracy: 0.9585
 Epoch 105/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3444 -
 accuracy: 0.9526 - val_loss: 0.3143 - val_accuracy: 0.9594
 Epoch 106/250
 675/675 [=====] - 10s 14ms/step - loss: 0.3429 -
 accuracy: 0.9518 - val_loss: 0.3125 - val_accuracy: 0.9606
 Epoch 107/250
 675/675 [=====] - 7s 11ms/step - loss: 0.3419 -
 accuracy: 0.9532 - val_loss: 0.3128 - val_accuracy: 0.9597
 Epoch 108/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3428 -
 accuracy: 0.9529 - val_loss: 0.3128 - val_accuracy: 0.9589
 Epoch 109/250
 675/675 [=====] - 7s 11ms/step - loss: 0.3414 -

accuracy: 0.9534 - val_loss: 0.3110 - val_accuracy: 0.9603
 Epoch 110/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3413 -
 accuracy: 0.9532 - val_loss: 0.3104 - val_accuracy: 0.9597
 Epoch 111/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3391 -
 accuracy: 0.9539 - val_loss: 0.3102 - val_accuracy: 0.9601
 Epoch 112/250
 675/675 [=====] - 10s 15ms/step - loss: 0.3399 -
 accuracy: 0.9525 - val_loss: 0.3104 - val_accuracy: 0.9604
 Epoch 113/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3386 -
 accuracy: 0.9535 - val_loss: 0.3089 - val_accuracy: 0.9606
 Epoch 114/250
 675/675 [=====] - 7s 11ms/step - loss: 0.3385 -
 accuracy: 0.9529 - val_loss: 0.3096 - val_accuracy: 0.9595
 Epoch 115/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3384 -
 accuracy: 0.9545 - val_loss: 0.3109 - val_accuracy: 0.9584
 Epoch 116/250
 675/675 [=====] - 7s 11ms/step - loss: 0.3385 -
 accuracy: 0.9529 - val_loss: 0.3090 - val_accuracy: 0.9604
 Epoch 117/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3366 -
 accuracy: 0.9542 - val_loss: 0.3087 - val_accuracy: 0.9600
 Epoch 118/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3360 -
 accuracy: 0.9529 - val_loss: 0.3072 - val_accuracy: 0.9592
 Epoch 119/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3358 -
 accuracy: 0.9546 - val_loss: 0.3063 - val_accuracy: 0.9606
 Epoch 120/250
 675/675 [=====] - 10s 14ms/step - loss: 0.3346 -
 accuracy: 0.9540 - val_loss: 0.3091 - val_accuracy: 0.9580
 Epoch 121/250
 675/675 [=====] - 8s 11ms/step - loss: 0.3357 -
 accuracy: 0.9532 - val_loss: 0.3067 - val_accuracy: 0.9598
 Epoch 122/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3353 -
 accuracy: 0.9541 - val_loss: 0.3058 - val_accuracy: 0.9604
 Epoch 123/250
 675/675 [=====] - 7s 11ms/step - loss: 0.3327 -
 accuracy: 0.9545 - val_loss: 0.3074 - val_accuracy: 0.9590
 Epoch 124/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3334 -
 accuracy: 0.9551 - val_loss: 0.3055 - val_accuracy: 0.9606
 Epoch 125/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3332 -

accuracy: 0.9536 - val_loss: 0.3042 - val_accuracy: 0.9603
 Epoch 126/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3331 -
 accuracy: 0.9538 - val_loss: 0.3047 - val_accuracy: 0.9604
 Epoch 127/250
 675/675 [=====] - 10s 14ms/step - loss: 0.3318 -
 accuracy: 0.9551 - val_loss: 0.3028 - val_accuracy: 0.9607
 Epoch 128/250
 675/675 [=====] - 8s 11ms/step - loss: 0.3307 -
 accuracy: 0.9551 - val_loss: 0.3022 - val_accuracy: 0.9613
 Epoch 129/250
 675/675 [=====] - 12s 18ms/step - loss: 0.3318 -
 accuracy: 0.9548 - val_loss: 0.3018 - val_accuracy: 0.9622
 Epoch 130/250
 675/675 [=====] - 13s 19ms/step - loss: 0.3294 -
 accuracy: 0.9556 - val_loss: 0.3028 - val_accuracy: 0.9613
 Epoch 131/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3295 -
 accuracy: 0.9556 - val_loss: 0.3011 - val_accuracy: 0.9603
 Epoch 132/250
 675/675 [=====] - 10s 15ms/step - loss: 0.3311 -
 accuracy: 0.9542 - val_loss: 0.3002 - val_accuracy: 0.9619
 Epoch 133/250
 675/675 [=====] - 12s 18ms/step - loss: 0.3308 -
 accuracy: 0.9545 - val_loss: 0.2998 - val_accuracy: 0.9621
 Epoch 134/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3290 -
 accuracy: 0.9550 - val_loss: 0.2993 - val_accuracy: 0.9624
 Epoch 135/250
 675/675 [=====] - 11s 17ms/step - loss: 0.3279 -
 accuracy: 0.9547 - val_loss: 0.2990 - val_accuracy: 0.9622
 Epoch 136/250
 675/675 [=====] - 11s 16ms/step - loss: 0.3286 -
 accuracy: 0.9550 - val_loss: 0.3001 - val_accuracy: 0.9618
 Epoch 137/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3276 -
 accuracy: 0.9552 - val_loss: 0.2987 - val_accuracy: 0.9616
 Epoch 138/250
 675/675 [=====] - 10s 15ms/step - loss: 0.3272 -
 accuracy: 0.9557 - val_loss: 0.2971 - val_accuracy: 0.9630
 Epoch 139/250
 675/675 [=====] - 10s 15ms/step - loss: 0.3276 -
 accuracy: 0.9553 - val_loss: 0.2993 - val_accuracy: 0.9613
 Epoch 140/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3265 -
 accuracy: 0.9553 - val_loss: 0.2970 - val_accuracy: 0.9628
 Epoch 141/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3245 -

accuracy: 0.9556 - val_loss: 0.2970 - val_accuracy: 0.9629
 Epoch 142/250
 675/675 [=====] - 7s 11ms/step - loss: 0.3248 -
 accuracy: 0.9564 - val_loss: 0.2968 - val_accuracy: 0.9625
 Epoch 143/250
 675/675 [=====] - 12s 17ms/step - loss: 0.3250 -
 accuracy: 0.9559 - val_loss: 0.2970 - val_accuracy: 0.9625
 Epoch 144/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3247 -
 accuracy: 0.9558 - val_loss: 0.2957 - val_accuracy: 0.9631
 Epoch 145/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3236 -
 accuracy: 0.9558 - val_loss: 0.2970 - val_accuracy: 0.9620
 Epoch 146/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3241 -
 accuracy: 0.9557 - val_loss: 0.2976 - val_accuracy: 0.9619
 Epoch 147/250
 675/675 [=====] - 7s 11ms/step - loss: 0.3234 -
 accuracy: 0.9557 - val_loss: 0.2950 - val_accuracy: 0.9627
 Epoch 148/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3225 -
 accuracy: 0.9553 - val_loss: 0.2947 - val_accuracy: 0.9629
 Epoch 149/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3237 -
 accuracy: 0.9551 - val_loss: 0.2937 - val_accuracy: 0.9622
 Epoch 150/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3219 -
 accuracy: 0.9557 - val_loss: 0.2935 - val_accuracy: 0.9620
 Epoch 151/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3206 -
 accuracy: 0.9567 - val_loss: 0.2928 - val_accuracy: 0.9624
 Epoch 152/250
 675/675 [=====] - 8s 13ms/step - loss: 0.3200 -
 accuracy: 0.9563 - val_loss: 0.2928 - val_accuracy: 0.9618
 Epoch 153/250
 675/675 [=====] - 10s 14ms/step - loss: 0.3193 -
 accuracy: 0.9566 - val_loss: 0.2921 - val_accuracy: 0.9633
 Epoch 154/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3202 -
 accuracy: 0.9563 - val_loss: 0.2955 - val_accuracy: 0.9609
 Epoch 155/250
 675/675 [=====] - 8s 13ms/step - loss: 0.3207 -
 accuracy: 0.9561 - val_loss: 0.2922 - val_accuracy: 0.9622
 Epoch 156/250
 675/675 [=====] - 14s 21ms/step - loss: 0.3197 -
 accuracy: 0.9568 - val_loss: 0.2911 - val_accuracy: 0.9631
 Epoch 157/250
 675/675 [=====] - 10s 15ms/step - loss: 0.3194 -

accuracy: 0.9567 - val_loss: 0.2903 - val_accuracy: 0.9631
 Epoch 158/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3183 -
 accuracy: 0.9571 - val_loss: 0.2923 - val_accuracy: 0.9618
 Epoch 159/250
 675/675 [=====] - 14s 20ms/step - loss: 0.3192 -
 accuracy: 0.9565 - val_loss: 0.2908 - val_accuracy: 0.9626
 Epoch 160/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3173 -
 accuracy: 0.9575 - val_loss: 0.2904 - val_accuracy: 0.9635
 Epoch 161/250
 675/675 [=====] - 10s 15ms/step - loss: 0.3176 -
 accuracy: 0.9577 - val_loss: 0.2902 - val_accuracy: 0.9624
 Epoch 162/250
 675/675 [=====] - 11s 17ms/step - loss: 0.3174 -
 accuracy: 0.9572 - val_loss: 0.2918 - val_accuracy: 0.9614
 Epoch 163/250
 675/675 [=====] - 12s 17ms/step - loss: 0.3162 -
 accuracy: 0.9575 - val_loss: 0.2893 - val_accuracy: 0.9634
 Epoch 164/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3166 -
 accuracy: 0.9567 - val_loss: 0.2895 - val_accuracy: 0.9624
 Epoch 165/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3184 -
 accuracy: 0.9565 - val_loss: 0.2923 - val_accuracy: 0.9616
 Epoch 166/250
 675/675 [=====] - 7s 11ms/step - loss: 0.3156 -
 accuracy: 0.9568 - val_loss: 0.2889 - val_accuracy: 0.9639
 Epoch 167/250
 675/675 [=====] - 11s 16ms/step - loss: 0.3167 -
 accuracy: 0.9569 - val_loss: 0.2884 - val_accuracy: 0.9631
 Epoch 168/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3147 -
 accuracy: 0.9575 - val_loss: 0.2889 - val_accuracy: 0.9637
 Epoch 169/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3165 -
 accuracy: 0.9559 - val_loss: 0.2875 - val_accuracy: 0.9635
 Epoch 170/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3146 -
 accuracy: 0.9563 - val_loss: 0.2883 - val_accuracy: 0.9630
 Epoch 171/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3148 -
 accuracy: 0.9567 - val_loss: 0.2889 - val_accuracy: 0.9634
 Epoch 172/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3136 -
 accuracy: 0.9573 - val_loss: 0.2879 - val_accuracy: 0.9634
 Epoch 173/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3146 -


accuracy: 0.9563 - val_loss: 0.2869 - val_accuracy: 0.9631
 Epoch 174/250
 675/675 [=====] - 10s 15ms/step - loss: 0.3125 -
 accuracy: 0.9576 - val_loss: 0.2871 - val_accuracy: 0.9639
 Epoch 175/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3124 -
 accuracy: 0.9574 - val_loss: 0.2858 - val_accuracy: 0.9650
 Epoch 176/250
 675/675 [=====] - 7s 11ms/step - loss: 0.3127 -
 accuracy: 0.9579 - val_loss: 0.2856 - val_accuracy: 0.9640
 Epoch 177/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3112 -
 accuracy: 0.9588 - val_loss: 0.2861 - val_accuracy: 0.9631
 Epoch 178/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3129 -
 accuracy: 0.9566 - val_loss: 0.2871 - val_accuracy: 0.9632
 Epoch 179/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3110 -
 accuracy: 0.9582 - val_loss: 0.2850 - val_accuracy: 0.9644
 Epoch 180/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3104 -
 accuracy: 0.9578 - val_loss: 0.2858 - val_accuracy: 0.9639
 Epoch 181/250
 675/675 [=====] - 8s 11ms/step - loss: 0.3101 -
 accuracy: 0.9577 - val_loss: 0.2838 - val_accuracy: 0.9636
 Epoch 182/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3099 -
 accuracy: 0.9579 - val_loss: 0.2851 - val_accuracy: 0.9633
 Epoch 183/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3086 -
 accuracy: 0.9576 - val_loss: 0.2819 - val_accuracy: 0.9643
 Epoch 184/250
 675/675 [=====] - 10s 15ms/step - loss: 0.3087 -
 accuracy: 0.9582 - val_loss: 0.2829 - val_accuracy: 0.9646
 Epoch 185/250
 675/675 [=====] - 11s 17ms/step - loss: 0.3093 -
 accuracy: 0.9576 - val_loss: 0.2828 - val_accuracy: 0.9640
 Epoch 186/250
 675/675 [=====] - 8s 11ms/step - loss: 0.3089 -
 accuracy: 0.9581 - val_loss: 0.2851 - val_accuracy: 0.9631
 Epoch 187/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3082 -
 accuracy: 0.9581 - val_loss: 0.2839 - val_accuracy: 0.9633
 Epoch 188/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3087 -
 accuracy: 0.9589 - val_loss: 0.2833 - val_accuracy: 0.9635
 Epoch 189/250
 675/675 [=====] - 10s 15ms/step - loss: 0.3063 -


accuracy: 0.9590 - val_loss: 0.2813 - val_accuracy: 0.9643
 Epoch 190/250
 675/675 [=====] - 11s 17ms/step - loss: 0.3075 -
 accuracy: 0.9591 - val_loss: 0.2805 - val_accuracy: 0.9643
 Epoch 191/250
 675/675 [=====] - 8s 11ms/step - loss: 0.3077 -
 accuracy: 0.9579 - val_loss: 0.2812 - val_accuracy: 0.9646
 Epoch 192/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3074 -
 accuracy: 0.9587 - val_loss: 0.2815 - val_accuracy: 0.9644
 Epoch 193/250
 675/675 [=====] - 8s 11ms/step - loss: 0.3055 -
 accuracy: 0.9595 - val_loss: 0.2802 - val_accuracy: 0.9648
 Epoch 194/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3071 -
 accuracy: 0.9581 - val_loss: 0.2805 - val_accuracy: 0.9633
 Epoch 195/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3068 -
 accuracy: 0.9584 - val_loss: 0.2797 - val_accuracy: 0.9645
 Epoch 196/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3057 -
 accuracy: 0.9588 - val_loss: 0.2790 - val_accuracy: 0.9655
 Epoch 197/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3063 -
 accuracy: 0.9594 - val_loss: 0.2794 - val_accuracy: 0.9639
 Epoch 198/250
 675/675 [=====] - 8s 11ms/step - loss: 0.3066 -
 accuracy: 0.9578 - val_loss: 0.2802 - val_accuracy: 0.9651
 Epoch 199/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3060 -
 accuracy: 0.9586 - val_loss: 0.2791 - val_accuracy: 0.9647
 Epoch 200/250
 675/675 [=====] - 7s 11ms/step - loss: 0.3061 -
 accuracy: 0.9583 - val_loss: 0.2787 - val_accuracy: 0.9646
 Epoch 201/250
 675/675 [=====] - 10s 14ms/step - loss: 0.3041 -
 accuracy: 0.9580 - val_loss: 0.2790 - val_accuracy: 0.9644
 Epoch 202/250
 675/675 [=====] - 8s 11ms/step - loss: 0.3060 -
 accuracy: 0.9573 - val_loss: 0.2783 - val_accuracy: 0.9650
 Epoch 203/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3046 -
 accuracy: 0.9590 - val_loss: 0.2788 - val_accuracy: 0.9638
 Epoch 204/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3049 -
 accuracy: 0.9580 - val_loss: 0.2800 - val_accuracy: 0.9639
 Epoch 205/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3021 -

accuracy: 0.9597 - val_loss: 0.2773 - val_accuracy: 0.9656
 Epoch 206/250
 675/675 [=====] - 10s 15ms/step - loss: 0.3037 -
 accuracy: 0.9588 - val_loss: 0.2770 - val_accuracy: 0.9642
 Epoch 207/250
 675/675 [=====] - 8s 11ms/step - loss: 0.3027 -
 accuracy: 0.9596 - val_loss: 0.2798 - val_accuracy: 0.9638
 Epoch 208/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3021 -
 accuracy: 0.9584 - val_loss: 0.2773 - val_accuracy: 0.9645
 Epoch 209/250
 675/675 [=====] - 9s 14ms/step - loss: 0.3021 -
 accuracy: 0.9590 - val_loss: 0.2768 - val_accuracy: 0.9648
 Epoch 210/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3026 -
 accuracy: 0.9589 - val_loss: 0.2755 - val_accuracy: 0.9648
 Epoch 211/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3023 -
 accuracy: 0.9592 - val_loss: 0.2772 - val_accuracy: 0.9644
 Epoch 212/250
 675/675 [=====] - 7s 11ms/step - loss: 0.3024 -
 accuracy: 0.9583 - val_loss: 0.2766 - val_accuracy: 0.9648
 Epoch 213/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3019 -
 accuracy: 0.9589 - val_loss: 0.2761 - val_accuracy: 0.9652
 Epoch 214/250
 675/675 [=====] - 8s 12ms/step - loss: 0.3025 -
 accuracy: 0.9596 - val_loss: 0.2763 - val_accuracy: 0.9658
 Epoch 215/250
 675/675 [=====] - 7s 11ms/step - loss: 0.2990 -
 accuracy: 0.9608 - val_loss: 0.2750 - val_accuracy: 0.9656
 Epoch 216/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3025 -
 accuracy: 0.9590 - val_loss: 0.2753 - val_accuracy: 0.9651
 Epoch 217/250
 675/675 [=====] - 7s 11ms/step - loss: 0.3005 -
 accuracy: 0.9600 - val_loss: 0.2767 - val_accuracy: 0.9648
 Epoch 218/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3018 -
 accuracy: 0.9586 - val_loss: 0.2742 - val_accuracy: 0.9647
 Epoch 219/250
 675/675 [=====] - 9s 13ms/step - loss: 0.2997 -
 accuracy: 0.9597 - val_loss: 0.2747 - val_accuracy: 0.9651
 Epoch 220/250
 675/675 [=====] - 8s 11ms/step - loss: 0.2990 -
 accuracy: 0.9606 - val_loss: 0.2748 - val_accuracy: 0.9650
 Epoch 221/250
 675/675 [=====] - 9s 13ms/step - loss: 0.3025 -

accuracy: 0.9575 - val_loss: 0.2756 - val_accuracy: 0.9644
 Epoch 222/250
 675/675 [=====] - 7s 11ms/step - loss: 0.2988 -
 accuracy: 0.9598 - val_loss: 0.2747 - val_accuracy: 0.9648
 Epoch 223/250
 675/675 [=====] - 9s 13ms/step - loss: 0.2988 -
 accuracy: 0.9591 - val_loss: 0.2744 - val_accuracy: 0.9651
 Epoch 224/250
 675/675 [=====] - 10s 15ms/step - loss: 0.2994 -
 accuracy: 0.9597 - val_loss: 0.2746 - val_accuracy: 0.9644
 Epoch 225/250
 675/675 [=====] - 7s 11ms/step - loss: 0.2996 -
 accuracy: 0.9590 - val_loss: 0.2750 - val_accuracy: 0.9651
 Epoch 226/250
 675/675 [=====] - 9s 13ms/step - loss: 0.2998 -
 accuracy: 0.9598 - val_loss: 0.2730 - val_accuracy: 0.9644
 Epoch 227/250
 675/675 [=====] - 8s 12ms/step - loss: 0.2975 -
 accuracy: 0.9603 - val_loss: 0.2763 - val_accuracy: 0.9641
 Epoch 228/250
 675/675 [=====] - 8s 12ms/step - loss: 0.2976 -
 accuracy: 0.9595 - val_loss: 0.2734 - val_accuracy: 0.9637
 Epoch 229/250
 675/675 [=====] - 9s 13ms/step - loss: 0.2968 -
 accuracy: 0.9602 - val_loss: 0.2732 - val_accuracy: 0.9647
 Epoch 230/250
 675/675 [=====] - 7s 11ms/step - loss: 0.2960 -
 accuracy: 0.9603 - val_loss: 0.2732 - val_accuracy: 0.9634
 Epoch 231/250
 675/675 [=====] - 9s 13ms/step - loss: 0.2979 -
 accuracy: 0.9594 - val_loss: 0.2717 - val_accuracy: 0.9656
 Epoch 232/250
 675/675 [=====] - 8s 12ms/step - loss: 0.2969 -
 accuracy: 0.9594 - val_loss: 0.2739 - val_accuracy: 0.9643
 Epoch 233/250
 675/675 [=====] - 9s 13ms/step - loss: 0.2963 -
 accuracy: 0.9597 - val_loss: 0.2731 - val_accuracy: 0.9640
 Epoch 234/250
 675/675 [=====] - 9s 13ms/step - loss: 0.2965 -
 accuracy: 0.9604 - val_loss: 0.2725 - val_accuracy: 0.9652
 Epoch 235/250
 675/675 [=====] - 7s 11ms/step - loss: 0.2978 -
 accuracy: 0.9596 - val_loss: 0.2731 - val_accuracy: 0.9641
 Epoch 236/250
 675/675 [=====] - 9s 13ms/step - loss: 0.2953 -
 accuracy: 0.9592 - val_loss: 0.2719 - val_accuracy: 0.9650
 Epoch 237/250
 675/675 [=====] - 8s 12ms/step - loss: 0.2962 -

accuracy: 0.9607 - val_loss: 0.2700 - val_accuracy: 0.9669
 Epoch 238/250
 675/675 [=====] - 8s 12ms/step - loss: 0.2943 -
 accuracy: 0.9597 - val_loss: 0.2708 - val_accuracy: 0.9659
 Epoch 239/250
 675/675 [=====] - 9s 13ms/step - loss: 0.2953 -
 accuracy: 0.9602 - val_loss: 0.2706 - val_accuracy: 0.9657
 Epoch 240/250
 675/675 [=====] - 8s 11ms/step - loss: 0.2953 -
 accuracy: 0.9596 - val_loss: 0.2698 - val_accuracy: 0.9660
 Epoch 241/250
 675/675 [=====] - 11s 16ms/step - loss: 0.2952 -
 accuracy: 0.9600 - val_loss: 0.2691 - val_accuracy: 0.9660
 Epoch 242/250
 675/675 [=====] - 7s 11ms/step - loss: 0.2941 -
 accuracy: 0.9599 - val_loss: 0.2693 - val_accuracy: 0.9666
 Epoch 243/250
 675/675 [=====] - 9s 13ms/step - loss: 0.2919 -
 accuracy: 0.9609 - val_loss: 0.2699 - val_accuracy: 0.9656
 Epoch 244/250
 675/675 [=====] - 9s 13ms/step - loss: 0.2916 -
 accuracy: 0.9611 - val_loss: 0.2693 - val_accuracy: 0.9657
 Epoch 245/250
 675/675 [=====] - 7s 11ms/step - loss: 0.2942 -
 accuracy: 0.9600 - val_loss: 0.2692 - val_accuracy: 0.9666
 Epoch 246/250
 675/675 [=====] - 9s 13ms/step - loss: 0.2939 -
 accuracy: 0.9594 - val_loss: 0.2698 - val_accuracy: 0.9656
 Epoch 247/250
 675/675 [=====] - 7s 11ms/step - loss: 0.2932 -
 accuracy: 0.9613 - val_loss: 0.2702 - val_accuracy: 0.9661
 Epoch 248/250
 675/675 [=====] - 9s 13ms/step - loss: 0.2922 -
 accuracy: 0.9606 - val_loss: 0.2689 - val_accuracy: 0.9663
 Epoch 249/250
 675/675 [=====] - 9s 13ms/step - loss: 0.2944 -
 accuracy: 0.9595 - val_loss: 0.2690 - val_accuracy: 0.9673
 Epoch 250/250
 675/675 [=====] - 8s 11ms/step - loss: 0.2923 -
 accuracy: 0.9609 - val_loss: 0.2672 - val_accuracy: 0.9668





Test Accuracy: 0.9688000082969666

In this experiment, we incorporated additional techniques to prevent overfitting and enhance the performance of the model. Firstly, we applied dropout layers with a dropout rate of 0.5 after both hidden layers. Dropout is a regularization technique that randomly sets a fraction of the input units to zero during training, which helps prevent overfitting by reducing interdependencies between neurons.

Additionally, we implemented early stopping based on validation loss values. This technique monitors the validation loss during training and stops the training process if the validation loss does not improve after a certain number of epochs (patience value of 10, in this case). Early stopping helps prevent overfitting by stopping the training process before the model starts to memorize the training data.

We also continued to apply L2 regularization to further control overfitting.

The results of these modifications are promising. The training accuracy achieved was 96.09%, indicating that the model performed well on the training set while preventing overfitting. The final test accuracy of 96.88% indicates that the model generalized well to unseen data and maintained

a high level of accuracy.

By analyzing the plot, we observed a significant decrease in the classification error until around 50 epochs, suggesting that the model quickly learned relevant patterns and improved its performance. After this point, the learning process gradually converged, resulting in a slower decrease in classification error.

In summary, the incorporation of dropout layers, early stopping, and L2 regularization helped control overfitting and improve the model's performance. The achieved accuracy on both training and test sets indicates the model's ability to generalize well and make accurate predictions on unseen data. The plot demonstrates the learning progression, with a significant reduction in classification error followed by a gradual convergence.

5 Que 5:

Try CNN (convolutional neural networks) on MNIST (or CIFAR or any dataset of your choice). Use the basic minibatch SGD as your learning algorithm, with or without regularizations. You may need to read Chapter-7, Chapter-8 and Chapter-9.

```
[22]: from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense

# Preprocess the dataset
x_train = train_images.reshape(-1, 28, 28, 1)
x_test = test_images.reshape(-1, 28, 28, 1)

# Convert labels to categorical
y_train = tf.keras.utils.to_categorical(train_labels, num_classes=10)
y_test = tf.keras.utils.to_categorical(test_labels, num_classes=10)

# Define the CNN architecture
model = Sequential()
model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model.add(MaxPooling2D((2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D((2, 2)))
model.add(Flatten())
model.add(Dense(64, activation='relu'))
model.add(Dense(10, activation='sigmoid'))

# Compile the model
model.compile(loss='categorical_crossentropy', optimizer=SGD(learning_rate=0.
↪01), metrics=['accuracy'])

# Train the model
history = model.fit(x_train, y_train, batch_size=64, epochs=11,
↪validation_split=0.2, verbose=1)

# Evaluate the model on the testing set
```


```
_, test_accuracy = model.evaluate(x_test, y_test, verbose=0)

print("\n")
# Plot the accuracy on training and validation data
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

print("\n\n")
# Print the test accuracy
print(f'Test Accuracy: {test_accuracy}')
```

```
Epoch 1/11
750/750 [=====] - 47s 62ms/step - loss: 1.0832 -
accuracy: 0.6881 - val_loss: 0.3104 - val_accuracy: 0.9128
Epoch 2/11
750/750 [=====] - 45s 60ms/step - loss: 0.2758 -
accuracy: 0.9168 - val_loss: 0.2009 - val_accuracy: 0.9416
Epoch 3/11
750/750 [=====] - 43s 58ms/step - loss: 0.1895 -
accuracy: 0.9432 - val_loss: 0.1553 - val_accuracy: 0.9559
Epoch 4/11
750/750 [=====] - 46s 61ms/step - loss: 0.1455 -
accuracy: 0.9561 - val_loss: 0.1227 - val_accuracy: 0.9651
Epoch 5/11
750/750 [=====] - 45s 60ms/step - loss: 0.1197 -
accuracy: 0.9633 - val_loss: 0.1095 - val_accuracy: 0.9682
Epoch 6/11
750/750 [=====] - 43s 58ms/step - loss: 0.1037 -
accuracy: 0.9684 - val_loss: 0.1057 - val_accuracy: 0.9680
Epoch 7/11
750/750 [=====] - 45s 60ms/step - loss: 0.0923 -
accuracy: 0.9717 - val_loss: 0.0870 - val_accuracy: 0.9745
Epoch 8/11
750/750 [=====] - 48s 64ms/step - loss: 0.0836 -
accuracy: 0.9746 - val_loss: 0.0822 - val_accuracy: 0.9759
Epoch 9/11
750/750 [=====] - 46s 62ms/step - loss: 0.0768 -
accuracy: 0.9768 - val_loss: 0.0839 - val_accuracy: 0.9754
Epoch 10/11
750/750 [=====] - 47s 63ms/step - loss: 0.0705 -
accuracy: 0.9787 - val_loss: 0.0726 - val_accuracy: 0.9780
Epoch 11/11
```

750/750 [=====] - 48s 64ms/step - loss: 0.0656 -
accuracy: 0.9803 - val_loss: 0.0845 - val_accuracy: 0.9748



Test Accuracy: 0.9764000177383423

In this final experiment, we constructed a neural network model using Convolutional Neural Network (CNN) layers, a popular choice for image classification tasks. For comparison purposes, we maintained the same number of layers, 11 epochs, a batch size of 64, and a learning rate of 0.01.

The results of this experiment are highly promising. Within just 11 epochs, the CNN model achieved an accuracy of 97.6% on the test data and 98.03% on train dataset. This high accuracy demonstrates the effectiveness of the CNN architecture in capturing relevant image features and making accurate predictions.

Additionally, the validation accuracy closely aligns with the test accuracy, indicating a low chance of overfitting. When the validation accuracy closely matches the test accuracy, it suggests that the model generalizes well to unseen data and can make reliable predictions on new instances.

The utilization of CNN layers in image classification models allows for the extraction of spatial hierarchies and local patterns present in the images. This enables the model to effectively learn and distinguish features essential for accurate classification.

Overall, the achieved accuracy of 97.6% on the test data, coupled with the similar validation accuracy, demonstrates the robustness and generalization capability of the CNN model. These results highlight the suitability of CNNs for image classification tasks, providing accurate predictions while minimizing the risk of overfitting.

[22] :

